

Processes

Daniel Hagimont (INPT)

hagimont@enseeiht.fr

<http://hagimont.perso.enseeiht.fr>

Process

- What difference between a process and a program
 - A program is passive
 - Stored on disk as an executable file
 - e.g. /bin/ls
 - A process is active
 - Execute on a processor

```
hagimont@hagimont-pc:~$ ls -la /bin/ls
-rwxr-xr-x 1 root root 142144 sept.  5  2019 /bin/ls
```

```
hagimont@hagimont-pc:~$ which ls
/usr/bin/ls
hagimont@hagimont-pc:~$ ls
bigdata2  Documents  install   Public    Téléchargements
Bureau    eclipse-workspace  Modèles  shared    tmp
divers    Images     Musique   snap      Vidéos
hagimont@hagimont-pc:~$
```

Process


- A process is an instance of a running program
 - Eg : gcc, sh, firefox ...
 - Created by the system or by an application
 - Created by a parent process
 - Uniquely identified (PID)
- Correspond to two units :
 - Execution unit
 - Sequential control flow (execute a flow of instructions)
 - Addressing unit
 - Each process has its own address space
 - Isolation

Process

- Processes can run on one or multiple processors
 - Several processes on one CPU: concurrency
 - Several processes on several CPU: parallelism

Concurrent processes


- Multiple processes can increase CPU utilization
 - Overlap one process's computation while another waits



- Multiple processes can reduce latency
 - Running A then B requires 100 secs for B to complete



- Running A and B concurrently (with preemption) improves the average response time

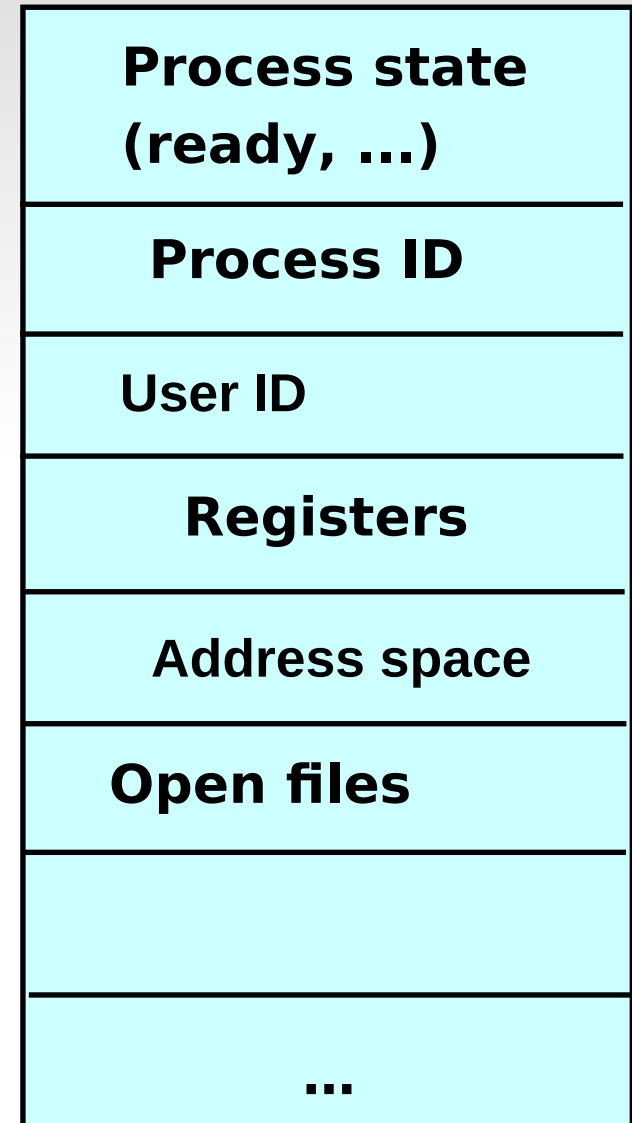


Execution context

- A process is characterized by its context
- Process' current state
 - Memory image
 - Code of the running program
 - Static and dynamic data
 - Register's state
 - Program counter (PC), Stack pointer (SP) ...
 - List of open files
 - Environment Variables
 - ...
- To be saved when the process is switched off
- To be restored when the process is switched on

Process Control Structure

- Hold a process execution context
- PCB (Process Control Block):
 - Data required by the OS to manage process
- Process tables:
 - PCB [MAX-PROCESSES]



Running mode

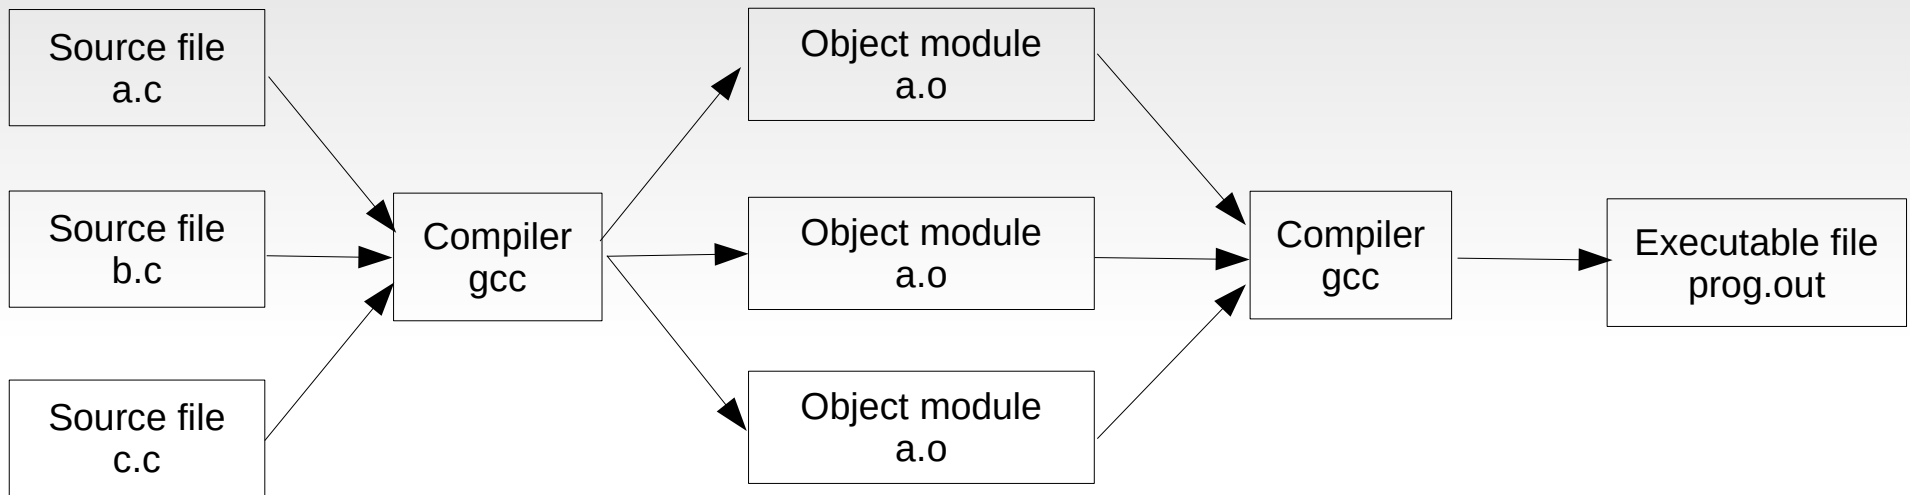
- User mode
 - Access restricted to process own address space
 - Limited instruction set
- Supervisor mode
 - Full memory access
 - Full access to the instruction set
- Interrupt, trap
 - Asynchronous event
 - Illegal instruction
 - System call request

Process memory layout

stack
free memory
heap
data
text


- Process execution state
 - Processor state
 - File descriptors
 - Memory allocation

Compiling



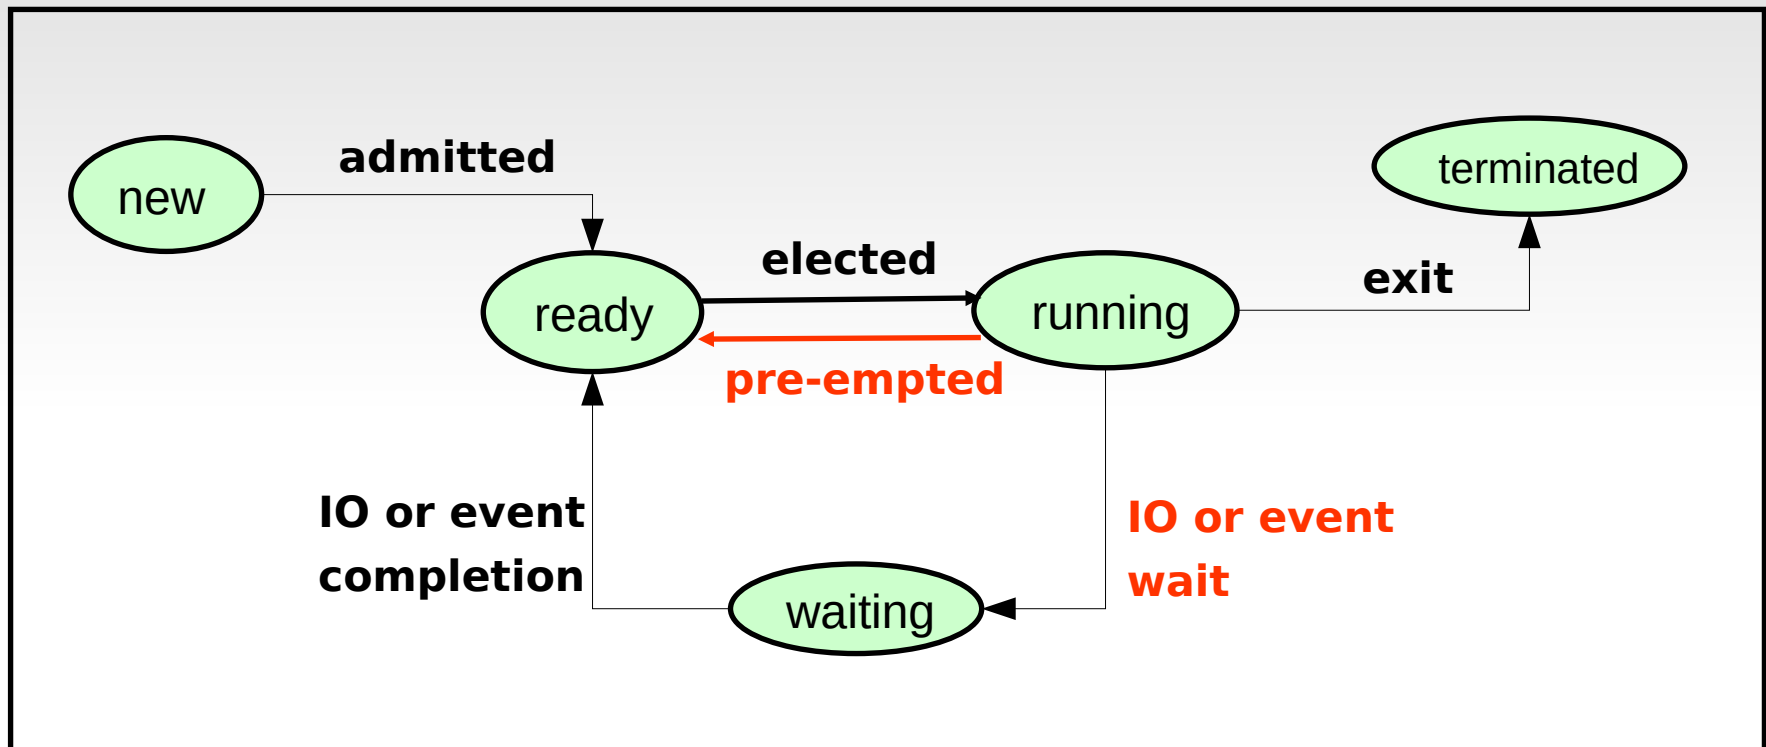
- Source files are compiled to object modules
- Object modules are linked into a single executable file
 - Example: `gcc <source> [-o output]`

Execute a process



- Create a new process (paused)
- Load executable file into process memory
- Load dynamic libraries
- Relocated APIs
- Set the program counter and stack pointer
- Resume the process

Process Lifecycle



- Which process should the kernel run ?
 - If 0 runnable, run a watchdog, if 1 runnable, run it
 - If n runnable, make scheduling decision

Exercise

- List all the running processes (with ps – see man)
- Start a new process (e.g. gnome-calculator)
- Find the id of this new process
- Show its status (see content of /proc/<id>/status)
- Pause it (kill with signal STOP)
- Resume it (kill with signal CONT)
- Terminate it (kill with signal KILL)
- Look at the tree of processes (pstree -a)

Process SVC overview

- `int fork ();`
 - Creates a new process that is an exact copy of the current one
 - Returns the process ID of the new process in the “parent”
 - Returns 0 in “child”
- `int waitpid (int pid, ...);`
 - pid – the process to wait for, or -1 for any
 - Returns pid of resuming process or -1 on error
- Hierarchy of processes
 - run the `ps tree -p` command

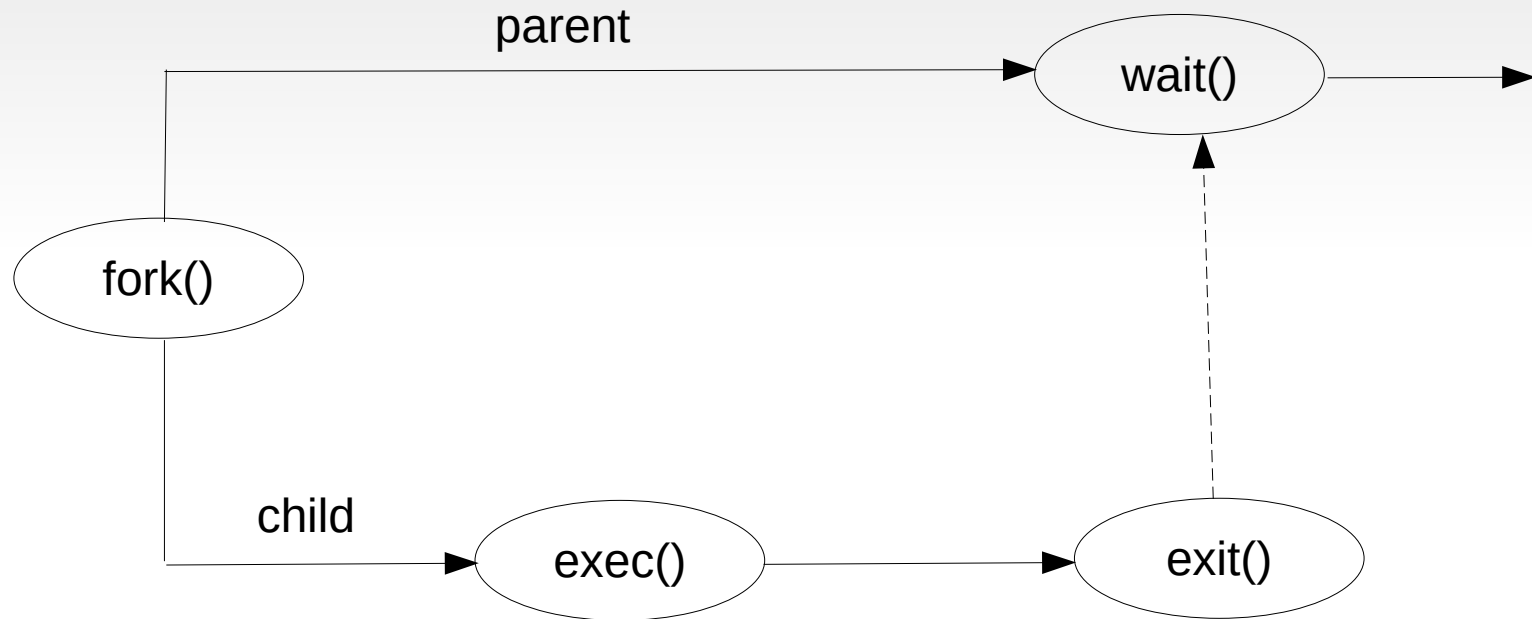
Process SVC overview

- `void exit (int status);`
 - Current process stops
 - status: returned to waitpid (shifted)
 - By convention, status of 0 is success
- `int kill (int pid, int sig);`
 - Sends signal sig to process pid
 - SIGTERM most common value, kills process by default (but application can catch it for “cleanup”)
 - SIGKILL stronger, kills process always
- When a parent process terminates before its child, 2 options:
 - Cascading termination (VMS)
 - Re-parent the orphan (UNIX)

Process SVC overview

- `int execve(const char *prog, const char **argv, char **envp;)`
 - prog – full pathname of program to run
 - argv – argument vector that gets passed to main
 - envp – environment variables, e.g., PATH, HOME
- Many other versions
 - `int execl(const char *path, const char *arg, ... /* (char *) NULL */);`
 - `int execlp(const char *file, const char *arg, ... /* (char *) NULL */);`
 - `int execlxe(const char *path, const char *arg, ... /*, (char *) NULL, char * const envp[] */);`
 - `int execv(const char *path, char *const argv[]);`
 - `int execvp(const char *file, char *const argv[]);`
 - `int execvpe(const char *file, char *const argv[], char *const envp[]);`

Process creation



Fork and Exec

- The fork system call creates a copy of the PCB
 - Opened files and memory mapped files are thus similar
 - Open files are thus opened by both father and child. They should both close the files.
 - The pages of many read only memory segments are shared (text, r/o data)
 - Many others are lazily copied (copy on write)
- The exec system call replaces the address space, the registers, the program counter by those of the program to exec.
 - But opened files are inherited

Why fork

- Most calls to fork followed by execvp
- Real win is simplicity of interface
 - Tons of things you might want to do to child
 - Fork requires no arguments at all
 - Without fork, require tons of different options
 - Example: Windows CreateProcess system call

```
Bool CreateProcess(  
    LPCTSTR lpApplicationName, //pointer to a name to executable module  
    LPTSTR lpCommandLine, // pointer to a command line string  
    LPSECURITY_ATTRIBUTES lpProcessAttributes, //process security attr  
    LPSECURITY_ATTRIBUTES lpThreadAttributes, // thread security attr  
    BOOL bInheritHandles, //creation flag  
    DWORD dwCreationFlags, // creation flags  
    LPVOID lpEnvironment, // pointer to new environment block  
    LPCTSTR lpCurrentDirectory, // pointer to current directory name  
    LPSTARTUPINFO lpStartupInfo, //pointer to STARTUPINFO  
    LPPROCESS_INFORMATION lpProcessInformation // pointer to PROCESS_INFORMATION );
```

Fork example

- Process creation
 - Done by cloning an existing process
 - Duplicate the process
 - Fork() system call
 - Return 0 to the child process
 - Return the child's pid to the father
 - Return -1 if error

```
#include <unistd.h>  
  
pid_t fork(void);
```

```
r = fork();  
if (r==-1) ... /* error */  
else if (r==0) ... /* child's code */  
else ... /* father's code */
```

Exercise

- How many processes are created ?

```
fork();  
fork();  
fork();
```

```
for (i=0; i<3;i++){  
    fork();  
}
```

- What are the possible different traces ?

```
int i = 0;  
switch (j=fork()) {  
    case -1 : perror("fork"); break;  
    case 0 : i++; printf("child :%d",i); break;  
    default : printf("father :%d",i);  
}
```

Exec example

- Reminder: main function definition
 - `int main(int argc, char *argv[]);`
- Execvp call
 - Replaces the process's memory image
 - `int execvp(const char *file, const char *argv[]);`
 - file : file name to load
 - argv : process parameters
 - execvp calls main(argc, argv) in the process to launch

Exercise

```
char * argv[3];  
argv[0] = "ls ";  
argv[1] = "-al ";  
argv[2] = NULL;  
execvp("ls", argv);
```

or

```
execvp("ls", "ls", "-al", NULL);
```

Father/child synchronization

- The father process waits for the completion of one of its child
 - `pid_t wait(int *status):`
 - The father waits for the completion of one of its child
 - `pid_t` : dead child's pid or -1 if no child
 - `status` : information on the child's death
 - `pid_t waitpid(pid_t pid, int *status, int option);`
 - Wait for a specific child's death
 - Option : non blocking ... see man

Wait example

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(){
    int spid, status;
    switch(spid = fork()){
        case -1 : perror("..."); exit(-1);
        case 0 : // child's code
            break;
        default : // the father wait for this child's terminaison
            if (waitpid(spid,&status,0)==-1) {perror("...");exit(-1);}
            ...
    }
}
```

Exercise (minishell)

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/wait.h>
#include <unistd.h>

pid_t pid;
char *av[2];
char cmd[20];

void doexec() {
    if (execvp(av[0],av)==-1)
        perror ("execvp failed");
    exit(0);
}
```

```
int main() {
    for (;;) {
        printf(">");
        scanf("%s",cmd);
        av[0] = cmd;
        av[1] = NULL;
        switch (pid = fork()) {
            case -1: perror("fork"); break;
            case 0:
                doexec();
            default:
                if (waitpid(pid, NULL, 0) == -1)
                    perror ("waitpid failed");
        }
    }
}
```

Resources you can read

- Operating System Concepts, 10th Edition, Abraham Silberschatz, Peter B. Galvin, Greg Gagne
 - <http://os-book.com/>
 - Chapters 3
- Modern Operating Systems, Andrew Tanenbaum
 - <http://www.cs.vu.nl/~ast/books/mos2/>
 - Chapter 2 (2.1)