

Inter Process Communication

Daniel Hagimont (INPT)

hagimont@enseeiht.fr

<http://hagimont.perso.enseeiht.fr>

I/O redirection

- A file is addressed through a descriptor
 - 0, 1 et 2 correspond to standard input, standard output, and standard error
 - The file descriptor number is returned by the open system call
- Basic operation
 - `int open(const char *pathname, int flags);`
 - `O_RDONLY`, `O_WRONLY`, `O_RDWR` ...
 - `int creat(const char *pathname, mode_t mode);`
 - `int close(int fd)`
 - `ssize_t read(int fd, void *buf, size_t count);`
 - `ssize_t write(int fd, void *buf, size_t count);`

I/O redirection

- Descriptor duplication
 - `dup(int oldfd); dup2(int oldfd, int newfd);`
 - Used to redirect standard I/O

```
#include <stdio.h>
#include <unistd.h>
int f;
/* redirect std input */

...
close(0);           // close std input
dup(f) ;            // duplicate f on the first free descriptor (i.e. 0)
close(f);           // free f
...
```

```
dup2(f,0);
close(f);
```

Exercise

- Copy with cat

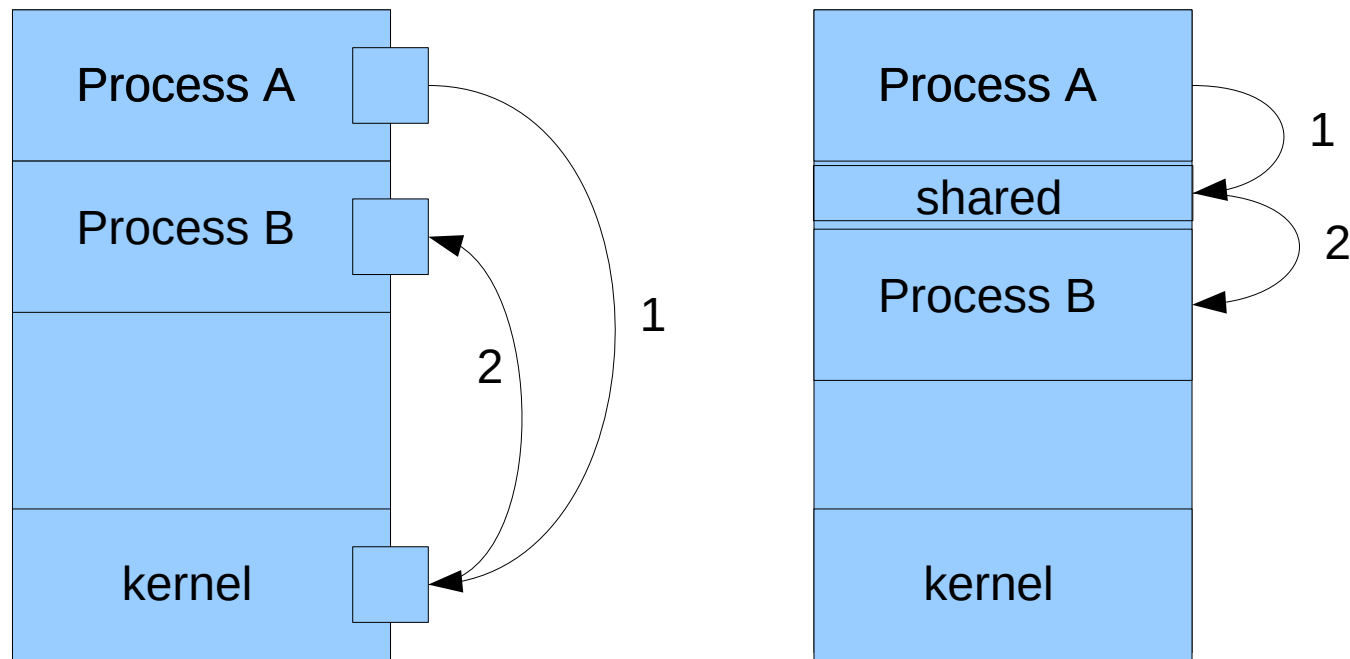
```
int main (int argc, char *argv[]) {
    char *src, *dest;
    int fd;
    if (argc != 3) {
        printf("usage: copy <src> <dest>\n");
        return 0;
    }
    src = argv[1];
    dest = argv[2];
    fd = open(dest, O_CREAT | O_WRONLY | O_TRUNC, S_IRUSR |
              S_IWUSR | S_IRGRP | S_IROTH);
    if (fd == -1) {perror("error open");exit(0);}
    if (dup2(fd,1) == -1) {perror("error dup2");exit(0);}
    if (execlp("cat","cat", src, NULL) == -1) {perror("error execl");exit(0);}
}
```

Cooperation between processes

- Independent process cannot affect or be affected by the execution of another process
- Cooperating process can affect or be affected by the execution of another process. Advantages:
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience

Process Interaction


- How can processes interact in real time?
 - Through files but it's not really “real time”.
 - Through asynchronous signals or alerts
 - By sharing a region of physical memory
 - By passing messages through the kernel/network



Pipe

- Communication mechanism between processes

- Fifo structure
- Limited capacity
- Producer/consumer synchronization



- `int pipe (int fds[2]);`


- Returns two file descriptors in `fds[0]` and `fds[1]`
- Writes to `fds[1]` will be read on `fds[0]`
- Returns 0 on success, -1 on error

- Operations on pipes

- read/write/close – as with files
- When `fds[1]` closed, `read(fds[0])` returns 0 bytes (EOF)
- When `fds[0]` closed, `write(fds[1])`: kill process with SIGPIPE

Exercise

```
int main (int argc, char *argv[]) {  
    int pipefds[2];  
    pipe (pipefds);  
    switch (fork ()) {  
        case -1:      perror ("fork"); exit (1);  
        case 0:  
            dup2 (pipefds[1], 1);  
            close (pipefds[0]); close (pipefds[1]);  
            execlp("ps", "ps", "-ef", NULL);  
  
        default:  
            dup2 (pipefds[0], 0);  
            close (pipefds[0]); close (pipefds[1]);  
            execlp("grep", "grep", "firefox", NULL);  
    }  
}
```



Asynchronous notification (Signal)

- A process may send a SIGSTOP, SIGTERM, SIGKILL signal to suspend (CTRL-Z), terminate or kill a process using the kill function:
 - `int kill (int pid, int sig);`
 - A lot of signals ... see man pages
 - Some signals cannot be blocked (SIGSTOP and SIGKILL)
- Upon reception of a signal, a given handler is called. This handler can be obtained and modified using the signal function:
 - `typedef void (*sighandler_t)(int); // handler`
 - `sighandler_t signal(int signum, sighandler_t handler); // set a handler`

Signal example

```
#include <stdio.h>
#include <stdlib.h>
#include <signal.h>
#include <unistd.h>

void handler(int signal_num) {
    printf("Signal %d => ", signal_num);
    switch (signal_num) {
        case SIGTSTP:
            printf("pause\n");
            break;
        case SIGINT:
        case SIGTERM:
            printf("End of the program\n");
            exit(0);
            break;
    }
}
```

```
int main(void) {
    signal(SIGTSTP, handler);
    /* if control-Z */
    signal(SIGINT, handler);
    /* if control-C */
    signal(SIGTERM, handler);
    /* if kill process */
    while (1) {
        sleep(1);
        printf(".\n");
    }
    printf("end");
    exit(0);
}
```

- Signal handling is vulnerable to race conditions: another signal (even of the same type) can be delivered to the process during execution of the signal handling routine.
- The `sigprocmask()` call can be used to block and unblock delivery of signals.

Exercise

- Without signals
 - Try control-C, control-Z
- With signals (previous slide)
 - Try control-C, control-Z

```
int main(void) {  
    while (1) {  
        sleep(1);  
        printf(".\n");  
    }  
}
```

Message queue

- Creation of a message queue
 - `int msgget(key_t key, int msgflg);`
- Control of the message queue
 - `int msgctl(int msqid, int cmd, struct msqid_ds *buf);`
- Emission of a message
 - `int msgsnd(int msqid, const void *msgp, size_t msgsz, int msgflg);`
- Reception of a message
 - `int msgrcv(int msqid, void *msgp, size_t msgsz, long msgtyp, int msgflg);`

Exercise

creator

```
int main() {
    int msgid;
    key_t key = 1234;

    /* Create the queue */
    if ((msgid = msgget(key,
        IPC_CREAT | 0666)) < 0) {
        perror("msgget failed");
        exit(1);
    }
}
```

sender

```
struct message {
    long mtype;
    char mtext[20];
};

int main() {
    int msgid;
    key_t key = 1234;
    struct message msg;

    /* get the queue */
    if ((msgid = msgget(key, 0666)) < 0) {
        perror("msgget failed");
        exit(1);
    }

    /* send a message */
    msg.mtype=1;
    strcpy(msg.mtext, "hello vietnam");
    if ((msgsnd(msgid, (void *)&msg,
        20,0)) == -1) {
        perror("msgsnd failed");
        exit(1);
    }
}
```

receiver

```
struct message {
    long mtype;
    char mtext[20];
};

int main() {
    int msgid;
    key_t key = 1234;
    struct message msg;

    /* get the queue */
    if ((msgid = msgget(key, 0666)) < 0) {
        perror("msgget failed");
        exit(1);
    }

    /* receive a message */
    if ((msgrcv(msgid, (void *)&msg,
        20,0,0)) == -1) {
        perror("msgsnd failed");
        exit(1);
    }
    printf("received : %s\n", msg.mtext);
}
```

Shared memory segment

- A process can create/use a shared memory segment using:
 - `int shmget(key_t key, size_t size, int shmflg);`
 - The returned value identifies the segment and is called the `shmid`
 - The key is used so that process indeed get the same segment.
- The owner of a shared memory segment can control access rights with `shmctl()`
- Once created, a shared segment should be attached to a process address space using
 - `void *shmat(int shmid, const void *shmaddr, int shmflg);`
- It can be detached using `int shmdt(const void *shmaddr);`
- Can also be done with the `mmap` function
- Example

Exercise

creator

```
int main() {
    int shmid;
    key_t key = 1234;
    /* Create the segment */
    if ((shmid = shmget(key, 10,
        IPC_CREAT | 0666)) < 0) {
        perror("shmget failed");
        exit(1);
    }
}
```

writer

```
int main() {
    int shmid, i, t;
    char *shm;
    key_t key = 1234;

    /* Get the segment */
    if ((shmid = shmget(key, 10,
        0666)) < 0) {
        perror("shmget failed");
        exit(1);
    }

    /* Attach the segment */
    if ((shm = shmat(shmid, NULL,
        0)) == (void *) -1) {
        perror("shmat failed");
        exit(1);
    }

    t = 0;
    while (1) {
        sleep(1);
        for (i=0; i<5; i++) shm[i] = 'a'+t;
        shm[i] = 0;
        printf("wrote : %s\n", shm);
        t++;
    }
}
```

reader

```
int main() {
    int shmid, i, t;
    char *shm;
    key_t key = 1234;

    /* Get the segment */
    if ((shmid = shmget(key, 10,
        0666)) < 0) {
        perror("shmget failed");
        exit(1);
    }

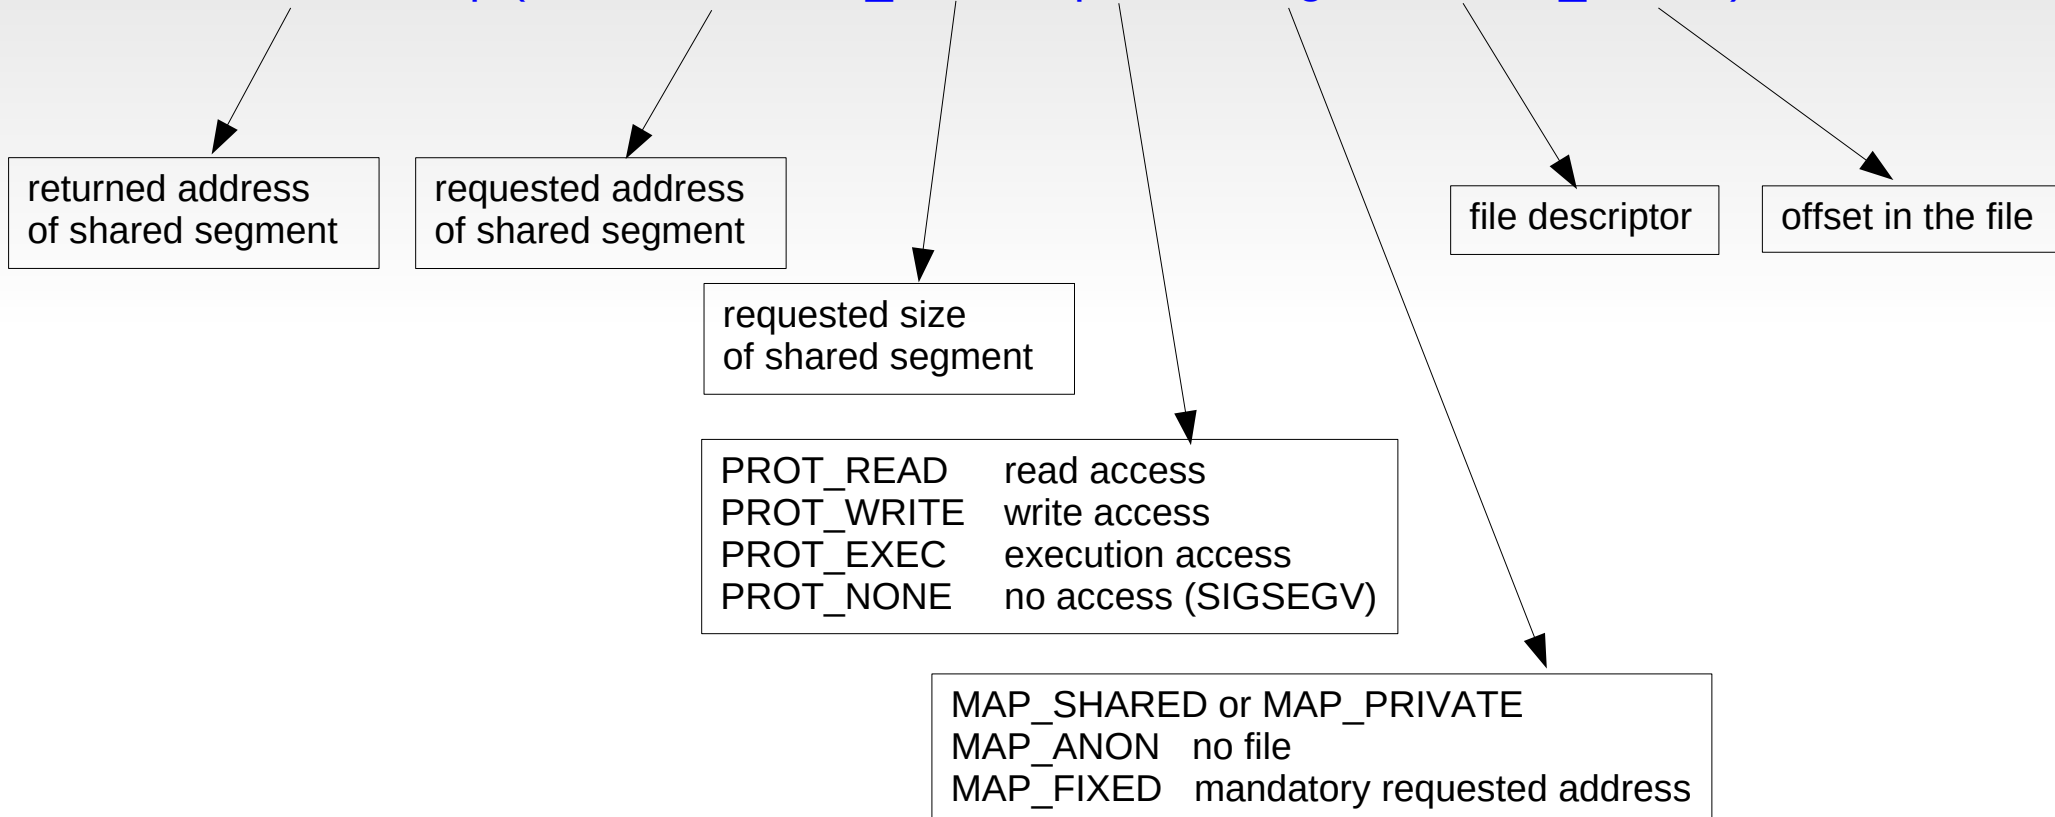
    /* Attach the segment */
    if ((shm = shmat(shmid, NULL,
        0)) == (void *) -1) {
        perror("shmat failed");
        exit(1);
    }

    while (1) {
        sleep(1);
        printf("read : %s\n", shm);
    }
}
```

Mmap

- Another interface for sharing memory

```
void * mmap (void * addr, size_t len, int prot, int flags, int fd, off_t offset);
```




Mmap examples

```
long pagesize = sysconf(_SC_PAGESIZE);  
int cf = open("content", O_RDWR);  
char* base = mmap(0, pagesize, PROT_WRITE|PROT_READ, MAP_SHARED, cf, 0);
```

```
char* b = mmap(0, pagesize, PROT_WRITE|PROT_READ,  
               MAP_SHARED|MAP_ANON, -1, 0);
```

/* addresses [base,base+pagesize[accessible in read/write mode */

Use-case: producer-consumer




```
#define BUFFER_SIZE 10
```

```
typedef struct {  
    char product;  
    int amount;  
} item;
```


```
item buffer [BUFFER_SIZE];  
int in = 0; // where to produce  
int out = 0; // where to consume  
int nb = 0; // number of items
```

Use-case: producer-consumer



```
void produce(item *i) {  
    while (nb == BUFFER_SIZE) {  
        // do nothing – no free place in buffer  
    }  
    memcpy(&buffer[in], i, sizeof(item));  
    in = (in+1) % BUFFER_SIZE;  
}
```

Use-case: producer-consumer



```
item *consume() {  
    item *i = malloc(sizeof(item));  
    while (nb == 0) {  
        // do nothing – nothing to consume  
    }  
    memcpy(i, &buffer[out], sizeof(item));  
    out = (out+1) % BUFFER_SIZE;  
    return i;  
}
```

Socket

- A socket is defined as an endpoint for communication
- Used for remote communication
- Basic message passing API
- Identified by an IP address and port
- The socket 161.25.19.8:1625 refers to port 1625 on host 161.25.19.8
- Communication between a pair of sockets and bidirectionnal
- => second part of Teaching Unit (networking)

Resources you can read

- Operating System Concepts, 10th Edition, Abraham Silberschatz, Peter B. Galvin, Greg Gagne
 - <http://os-book.com/>
 - Chapters 3
- Modern Operating Systems, Andrew Tanenbaum
 - <http://www.cs.vu.nl/~ast/books/mos2/>
 - Chapter 2 (2.3)