

Threads

Daniel Hagimont (INPT)

hagimont@enseeiht.fr

<http://hagimont.perso.enseeiht.fr>

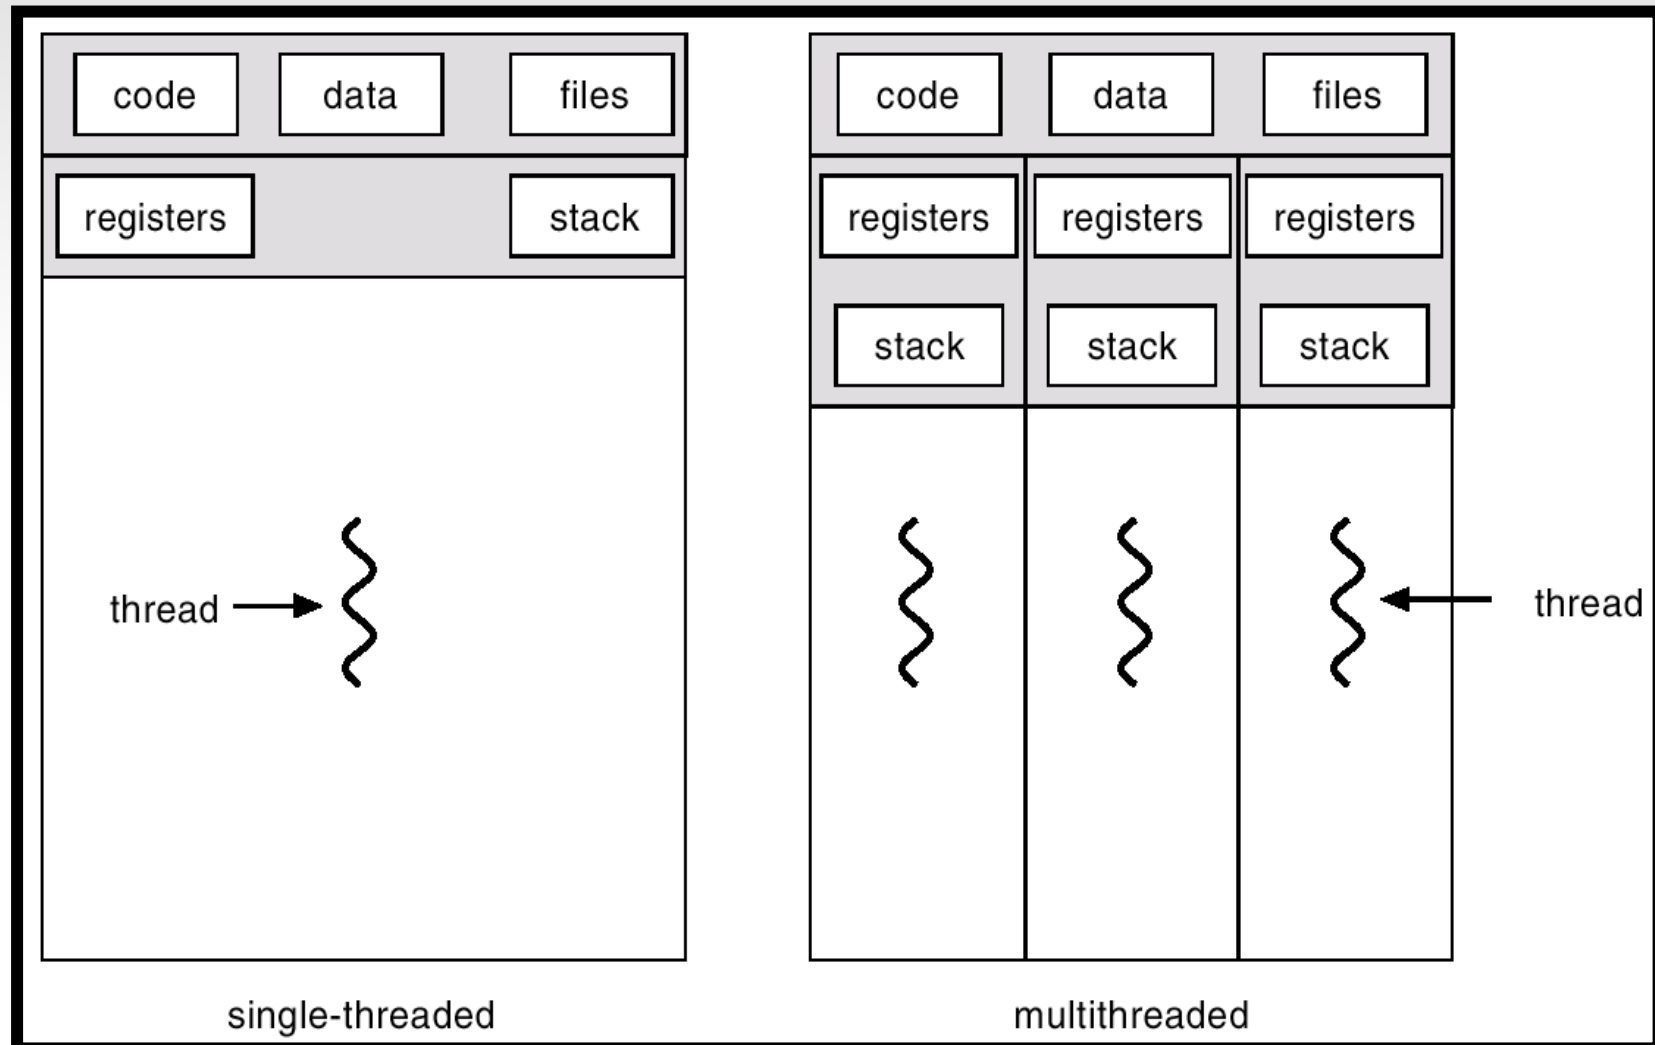
Process

- Unix process: heavy
 - Context: large data structure (includes an address space)
 - Protected address space
 - Address space not accessible from another process
 - Sharing / communication
 - At creation time (fork)
 - Via shared memory segments
 - Via messages (queues, sockets)
 - Communication is costly

Threads

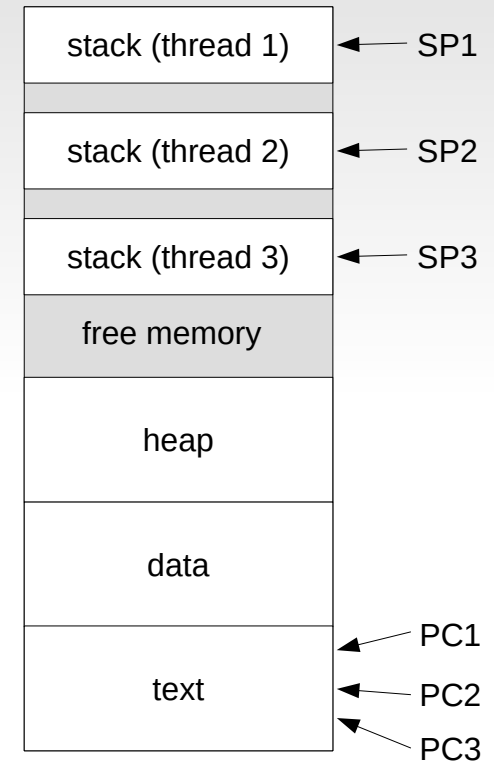
- Light weight process
 - Light weight context
 - A shared context: address space, open files ...
 - A private context: stack, registers ...
- Faster communication within the same address space
 - Message exchange, shared memory, synchronization
- Useful for concurrent/parallel applications
 - Easier
 - More efficient
 - Multi-core processors

Single-threaded vs multi-threaded processes



Multi-threaded process

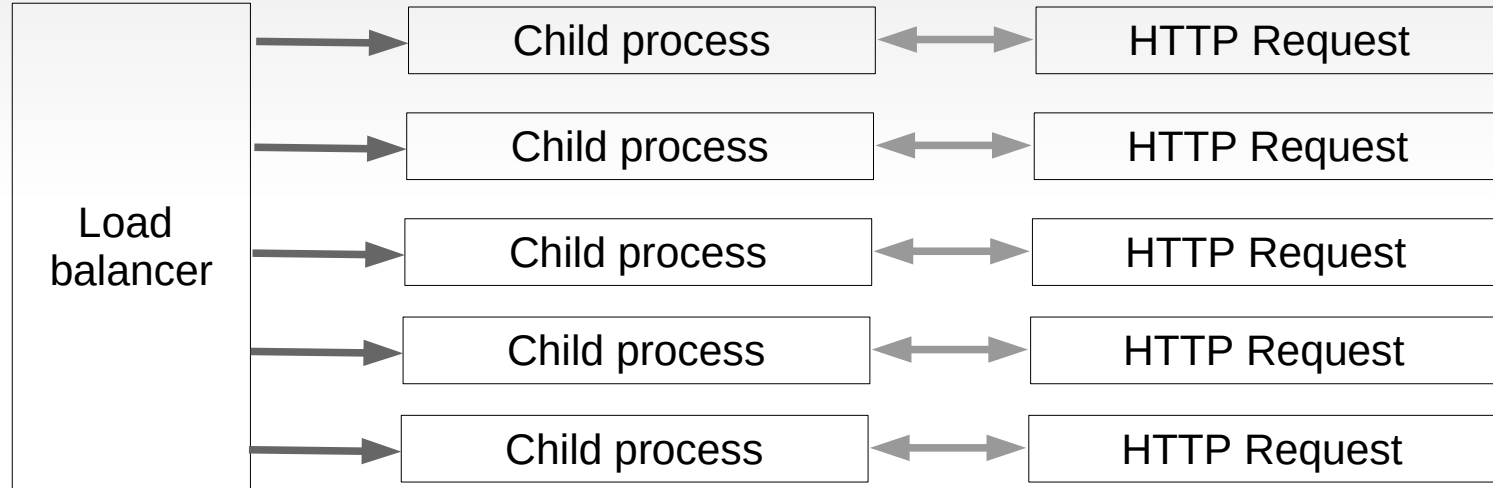
- Each thread has:
 - Private stack
 - Private stack pointer
 - Private program counter private register values
- Share:
 - Common text section (code)
 - Common data section (global data)
 - Common heap (dynamic allocations)
 - File descriptors (opened files)
 - Signals



Multi-threaded process

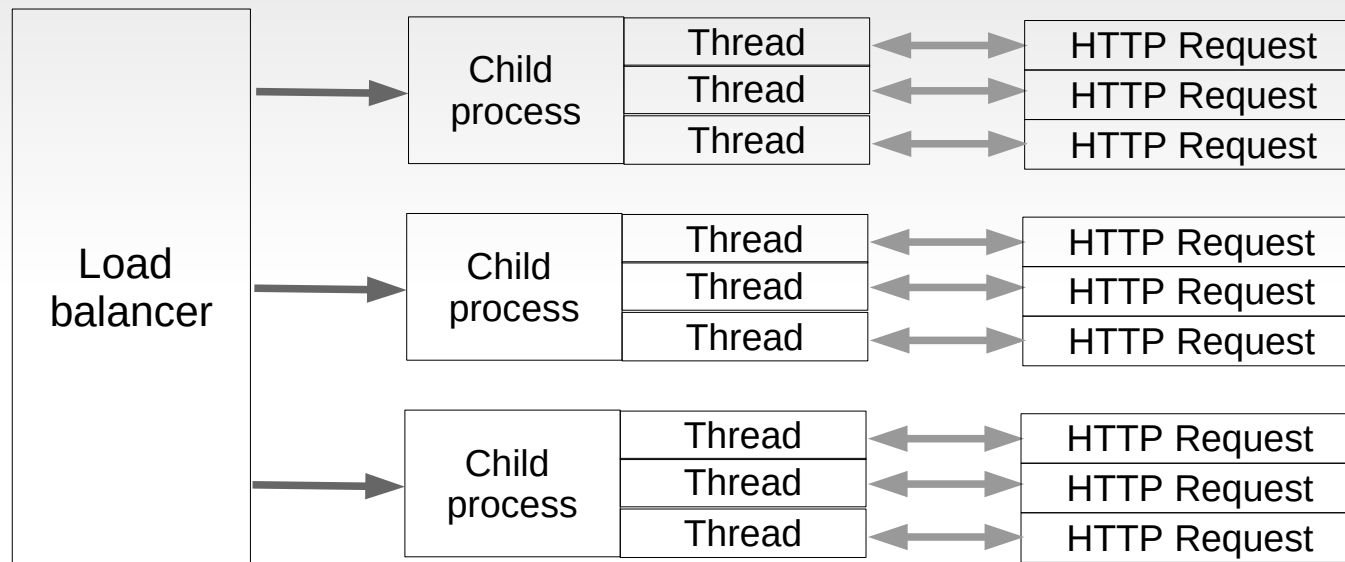
- Threads: same goal as processes
 - Do several thing at the same time
 - Increase CPU utilization
 - Increase responsiveness
- What's the difference
 - Multi-process with `fork()`: resource cloning
 - Multi-thread process: resource sharing

Some multi-process architectures



Apache HTTPD Prefork Model

Some multi-process architectures



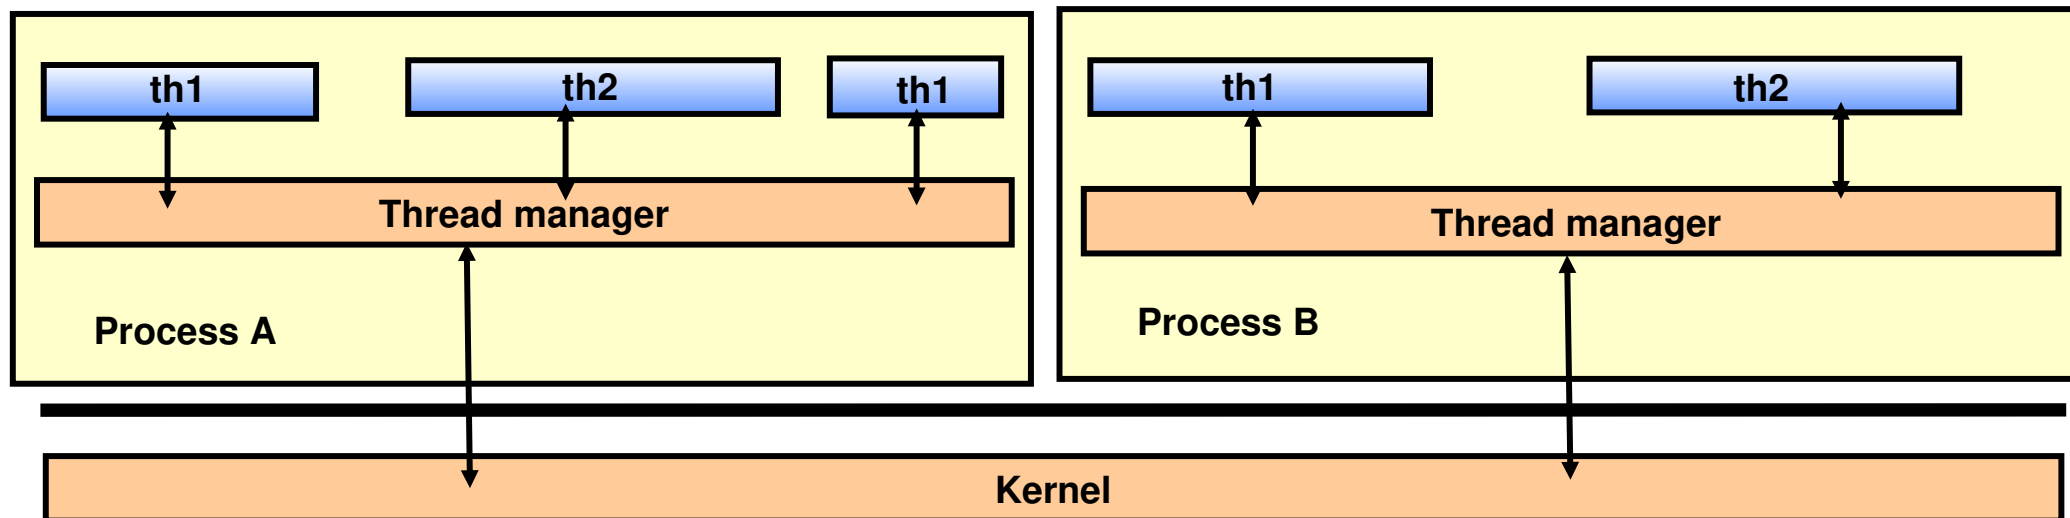
Apache HTTPD Worker Model

Exercise

- Show the number of threads for process firefox or google-chrome
 - ps with NLWP (number of lightweight processes) option
 - e.g. ps -o nlwp <processId>
 - Count number of subdirectories in /proc/<processId>/task

User-level Threads

- Implemented in a user level library
- Unmodified Kernel
- Threads and thread scheduler run in the same user process




Advantages and disadvantages of User-level threads

- Parallelism (-)
 - No real parallelism between the threads within a process
- Efficiency (+)
 - Quick context switch
- Blocking system call (-)
 - The process is blocked in the kernel
 - All thread are blocked until the system call (I/O) is terminated

Kernel level threads

- Thread managed by the kernel
- Thread creation as a system call
- When a thread is blocked, the processor is allocated to another thread by the kernel



Advantages and disadvantages of Kernel-level threads

- Blocking system call (+)
 - When a thread is blocked due to an SVC call, threads in the same process are not
- Real Parallelism (+)
 - N threads in the same process can run on K processors (multi-core)
- Efficiency (-)
 - More expensive context switch / user level threads
 - Every management operation goes through the kernel
 - Require more memory

POSIX Threads : Pthreads API

- `int pthread_create (pthread_t *thread, const pthread_attr_t *attr, void * (*start_routine)(void *), void *arg);`
 - Creates a thread
- `pthread_t pthread_self (void);`
 - Returns id of the current thread
- `int pthread_equal (pthread_t thr1, pthread_t thr2);`
 - Compare 2 thread ids
- `void pthread_exit (void *status);`
 - Terminates the current thread
- `int pthread_join (pthread_t thr, void **status);`
 - Waits for completion of a thread
- `int pthread_yield(void);`
 - Relinquish the processor
- Plus lots of support for synchronization [next lecture]

Exercise (pthread) (1/2)

```
#include <pthread.h>

void * ALL_IS_OK = (void *)123456789L;

char *mess[2] = { "thread1", "thread2" };

void * writer(void * arg)
{
    int i, j;

    for(i=0;i<10;i++) {
        printf("Hi %s! (I'm %lx)\n", (char *) arg, pthread_self());
        j = 800000; while(j!=0) j--;
    }

    return ALL_IS_OK;
}
```

Exercise (pthread) (2/2)

```
int main(void)
{ void * status;
  pthread_t writer1_pid, writer2_pid;

  pthread_create(&writer1_pid, NULL, writer, (void *)mess[1]);
  pthread_create(&writer2_pid, NULL, writer, (void *)mess[0]);

  pthread_join(writer1_pid, &status);
  if(status == ALL_IS_OK)
    printf("Thread %lx completed ok.\n", writer1_pid);

  pthread_join(writer2_pid, &status);
  if(status == ALL_IS_OK)
    printf("Thread %lx completed ok.\n", writer2_pid);

  return 0;
}
```


Fork(), exec()

- What happens if one thread of a program calls fork()?
 - Does the new process duplicate all threads ? Or is the newprocess single-threaded ?
 - Some UNIX systems have chosen to have two versions of fork()
- What happens if one thread of a program calls exec()?
 - Generally, the new program replace the entire process, including all threads.

Resources you can read

- Pthreads
 - <https://computing.llnl.gov/tutorials/pthreads/>
- Operating System Concepts, 10th Edition, Abraham Silberschatz, Peter B. Galvin, Greg Gagne
 - <http://os-book.com/>
 - Chapters 4
- Modern Operating Systems, Andrew Tanenbaum
 - <http://www.cs.vu.nl/~ast/books/mos2/>
 - Chapter 2 (2.2)