

Memory management

Daniel Hagimont (INPT)

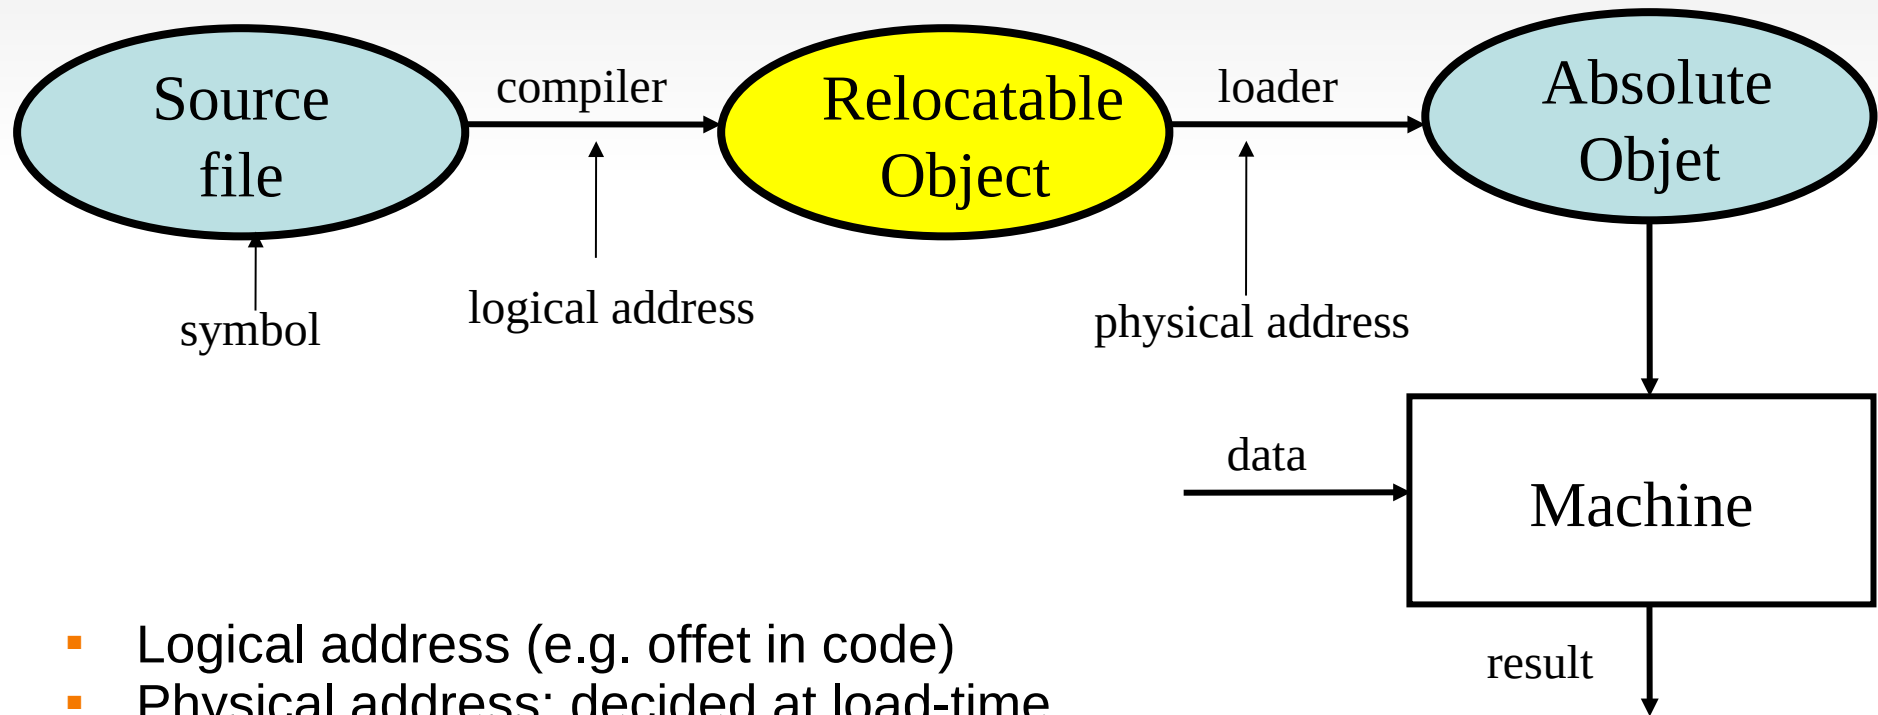
hagimont@enseeiht.fr

<http://hagimont.perso.enseeiht.fr>

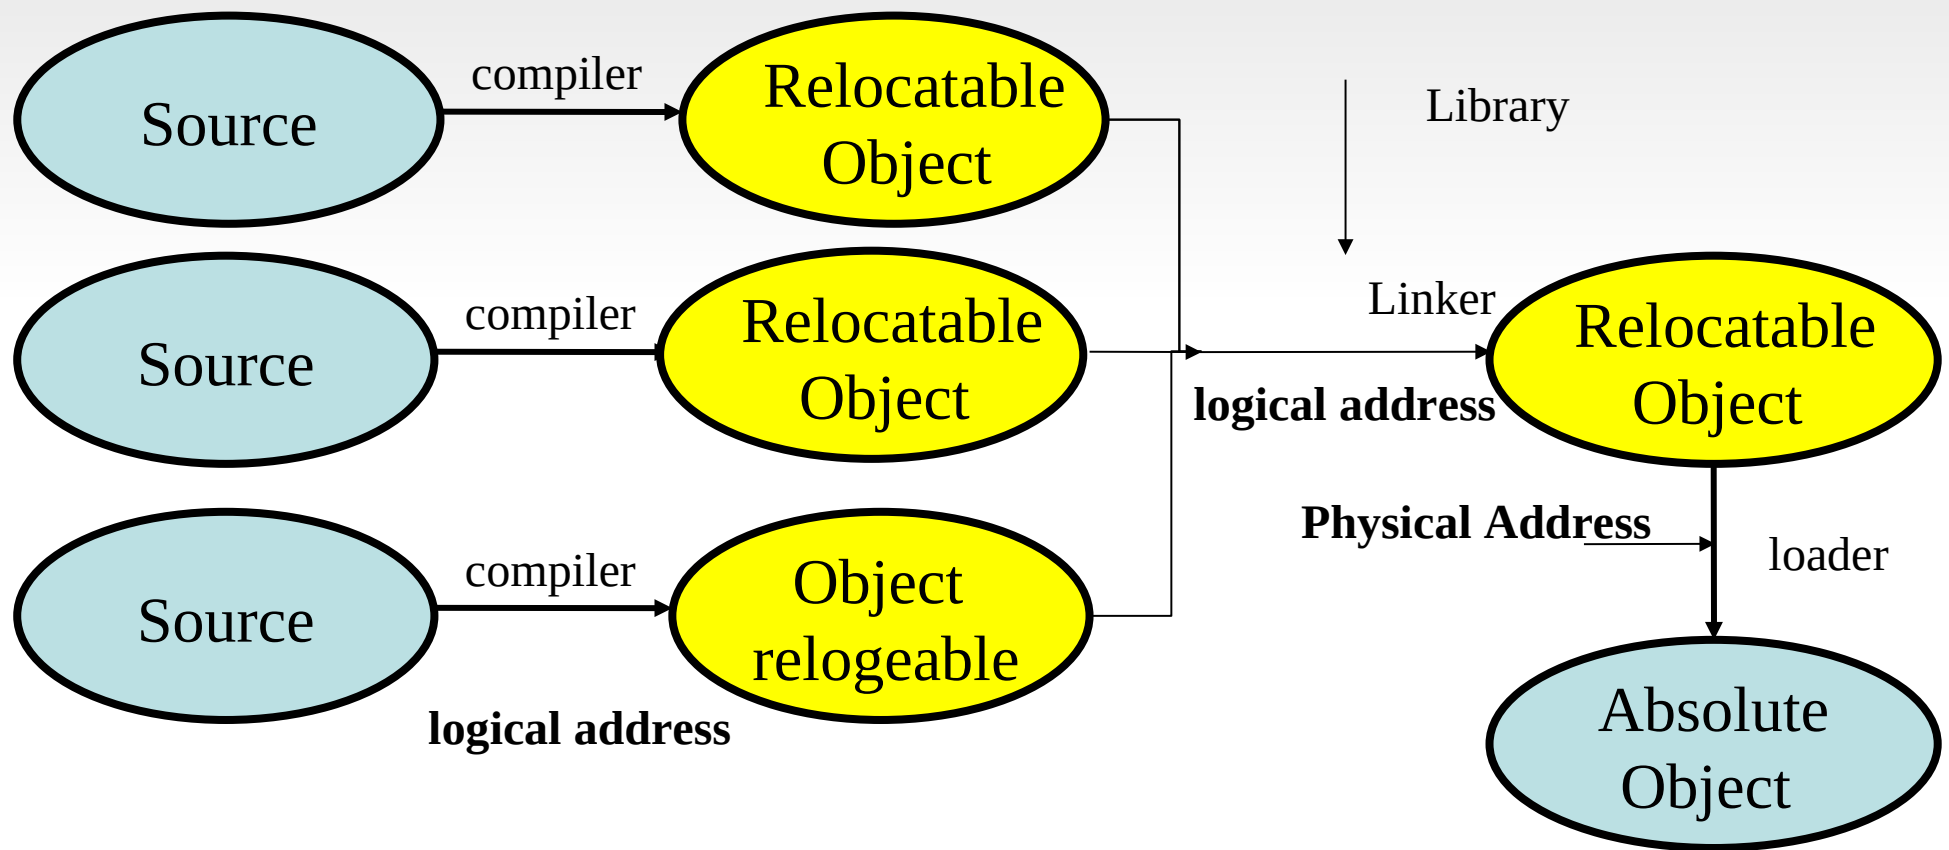
Introduction

- Memory is a resource required by all processes
 - Every program needs to be loaded in memory to be running
- Problems
 - Address translation
 - Symbol → Logical address → physical address
 - Memory allocation and exhaustion
 - Memory sharing
 - Memory protection

Life cycle of a single program



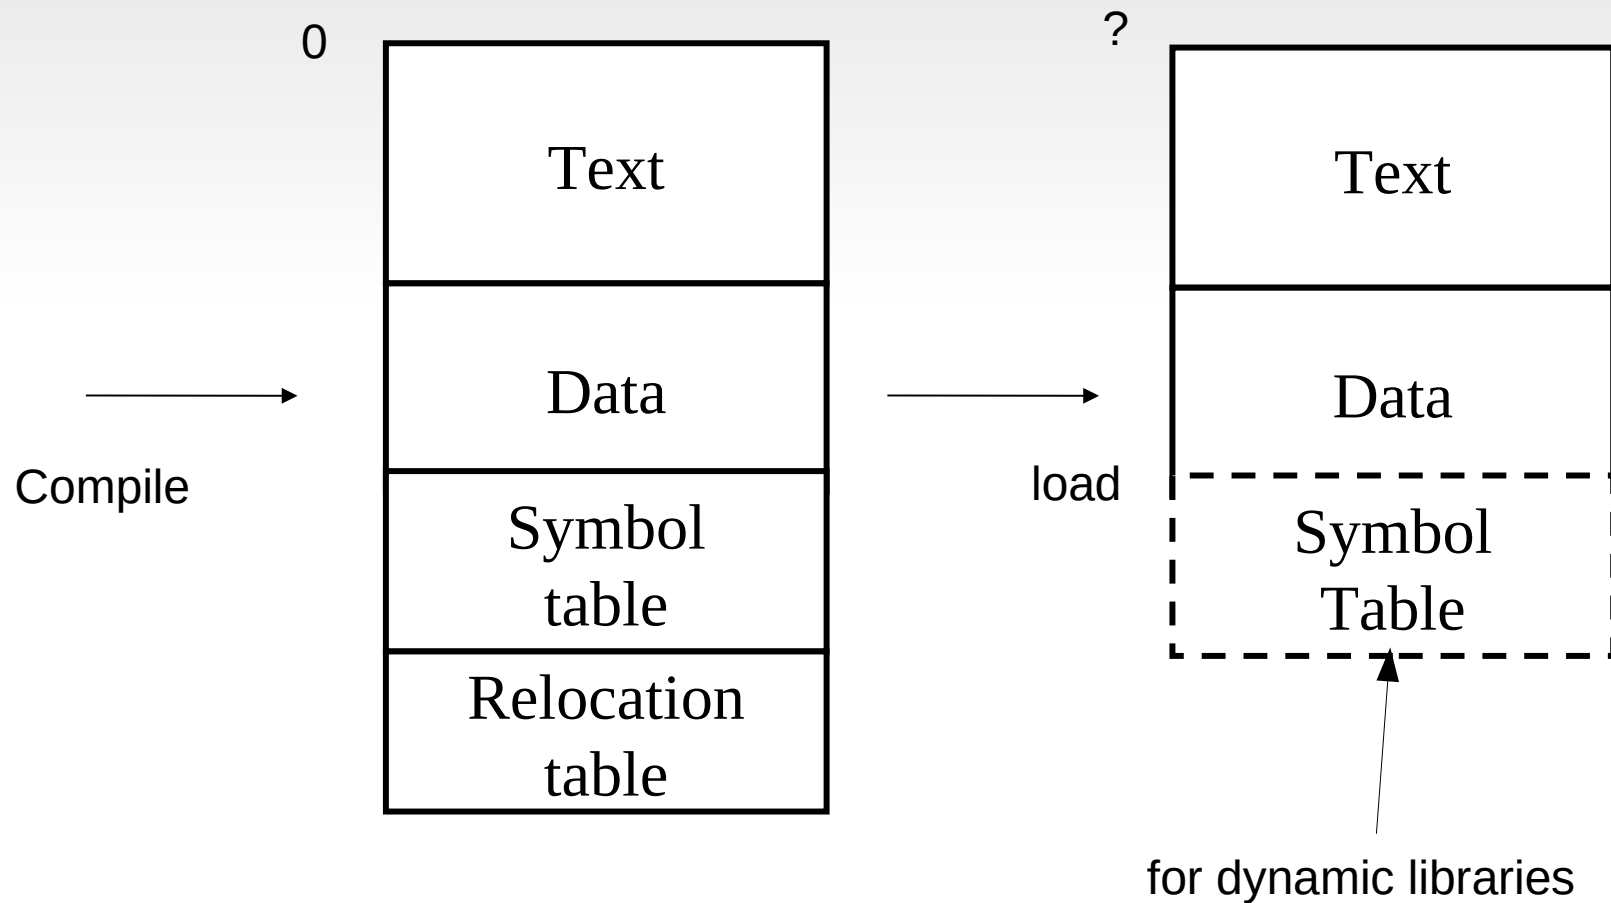
Lifecycle of a program assembled from multiple parts



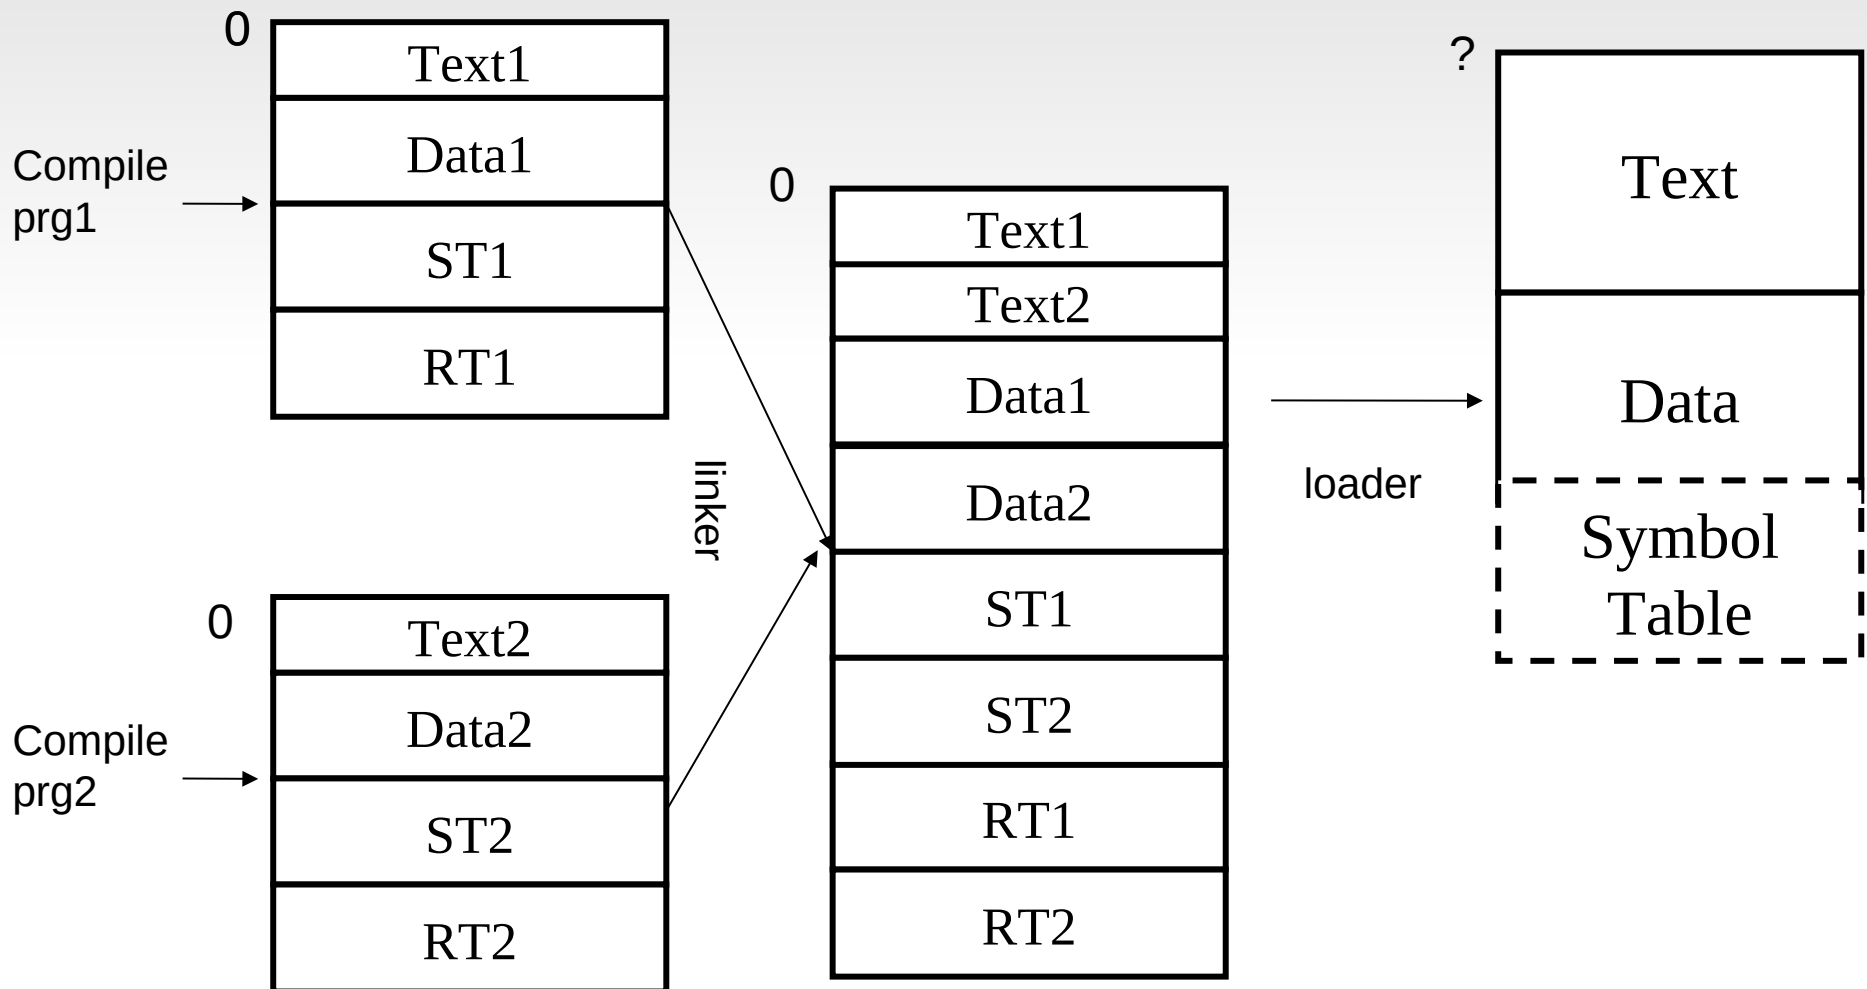
Load-time translation

- Translation between logical and physical addresses
 - Determine where process will reside in memory
 - Translate all references within program
 - Established once for all
- Monoprogramming
 - One program in memory
 - Easy (could even be done before load-time)
- Multiprogramming
 - N programs in memory
 - Compiler and linker do not know the implantation of processes in memory
 - Need to track op-codes that must be updated at load-time

Simple program binary structure



Complex program binary structure



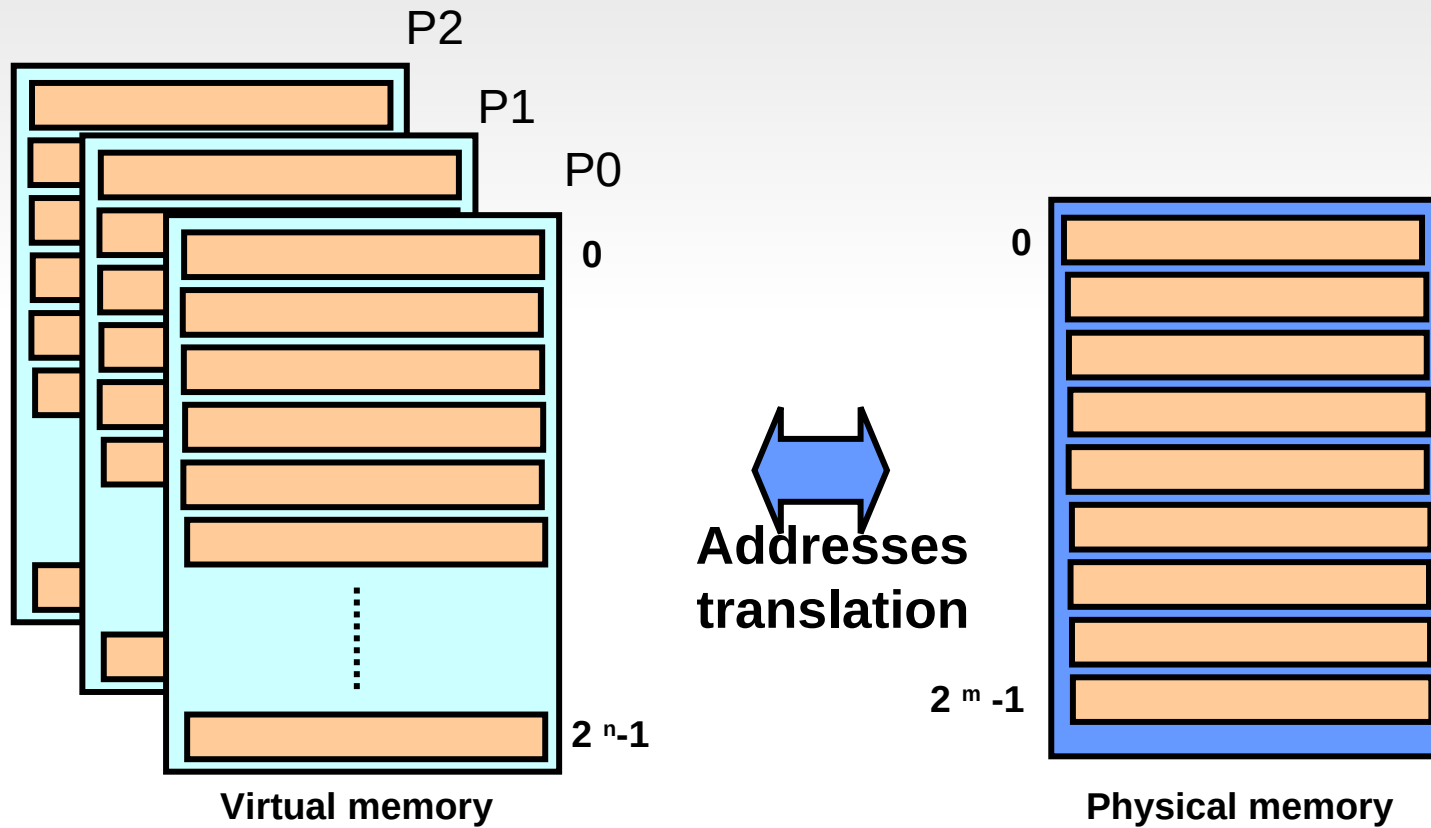
Load-time translation summary

- Remaining problems
 - How to enforce protection ?
 - How to move program once in memory ?
 - What if no contiguous free region fits program size ?
 - Can we separate linking from memory management problems ?

Virtual memory

- Separate linking problem from memory management
- Give each program its own virtual address space
 - Linker works on virtual addresses
 - Virtual address translation done at runtime
 - Relocate each load/store to its physical address
 - Require specific hardware (MMU)

Virtual memory



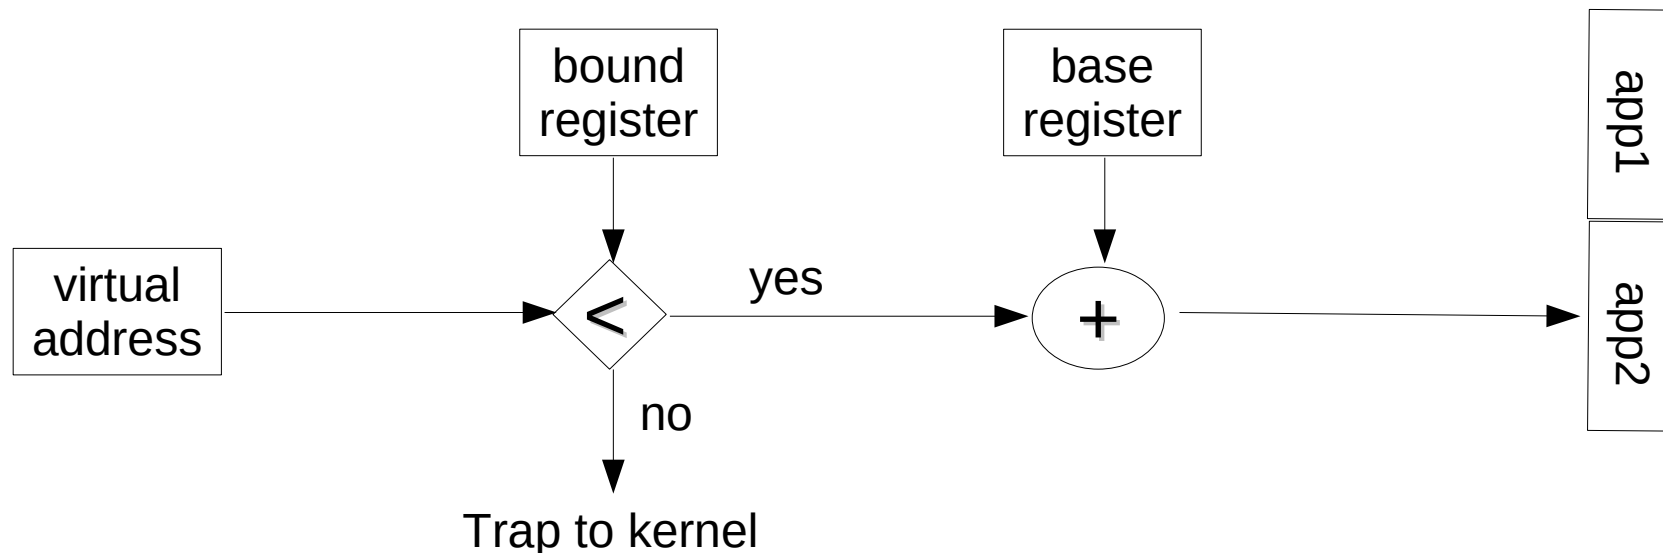
Ideally we want to enable $n > m$ and non contiguous allocation

Virtual memory expected benefits

- Programs can be relocated while running
 - Ease swap in/swap out
- Enforce protection
 - Prevent one app from messing with another's memory
- Programs can see more memory than exists
 - Most of a process's memory will be idle
 - Write idle part to disk until needed (swap)

1st idea : Base + bound registers

- Contiguous allocation of variable size
- Two special privileged registers: base and bound
- On each load/store:
 - Check $0 \leq \text{virtual address} < \text{bound}$, else trap to kernel
 - Physical address = virtual address (plus) base



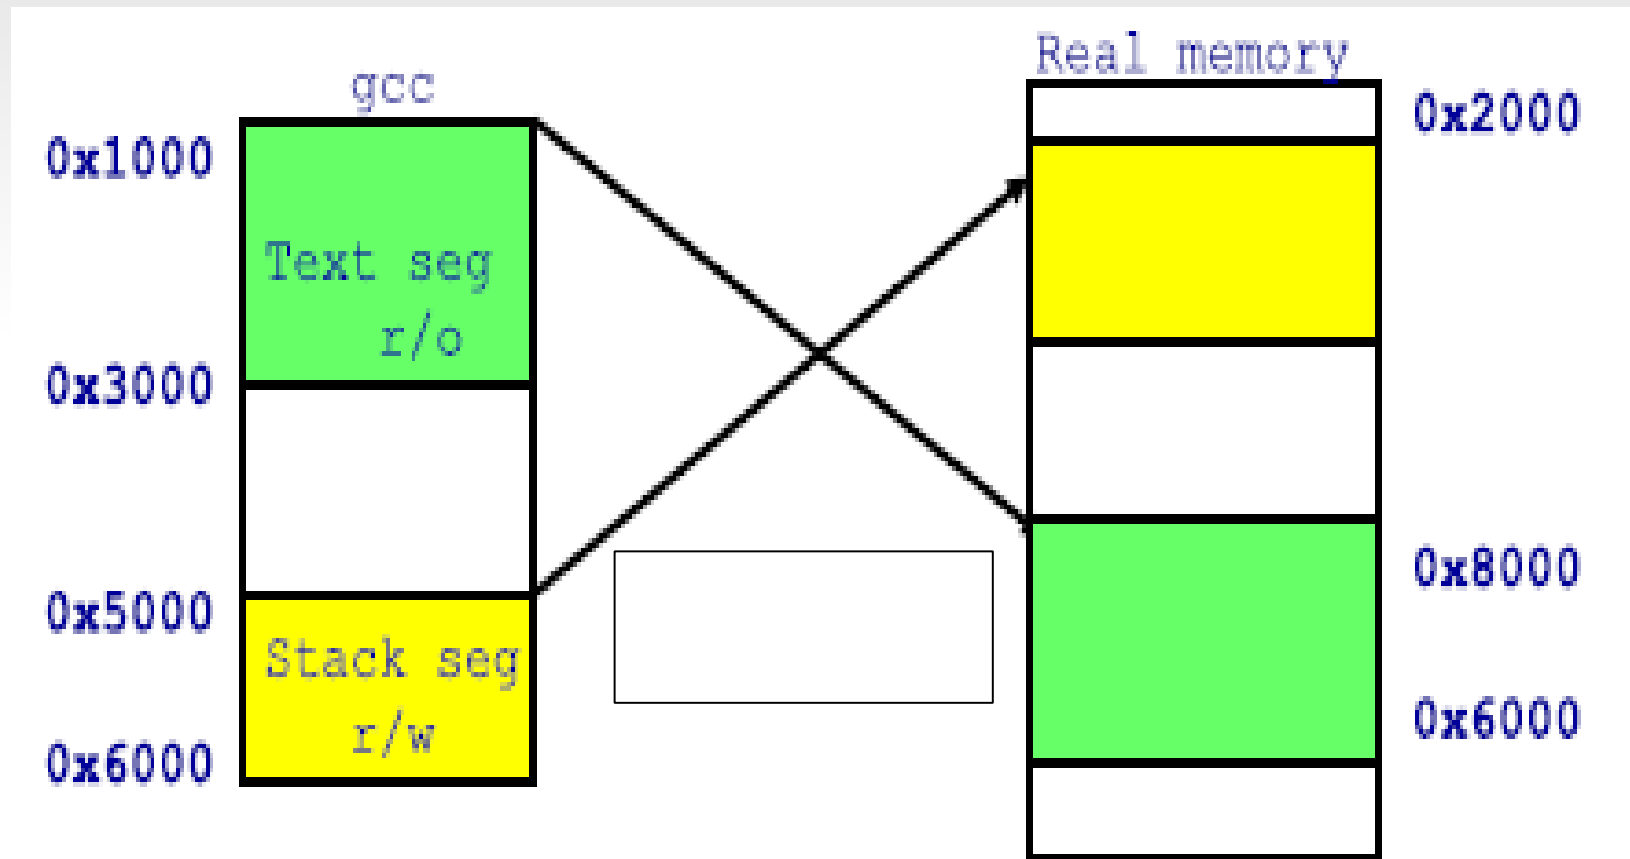
Base + bounds register

- Moving a process in memory
 - Change base register
- Context switch
 - OS must re-load base and bound register
- Advantages
 - Cheap in terms of hardware: only two registers
 - Cheap in terms of cycles: do add and compare in parallel
- Disadvantages
 - Still contiguous allocation
 - Growing a process is expensive or impossible
 - Hard to share code or data

Segmentation

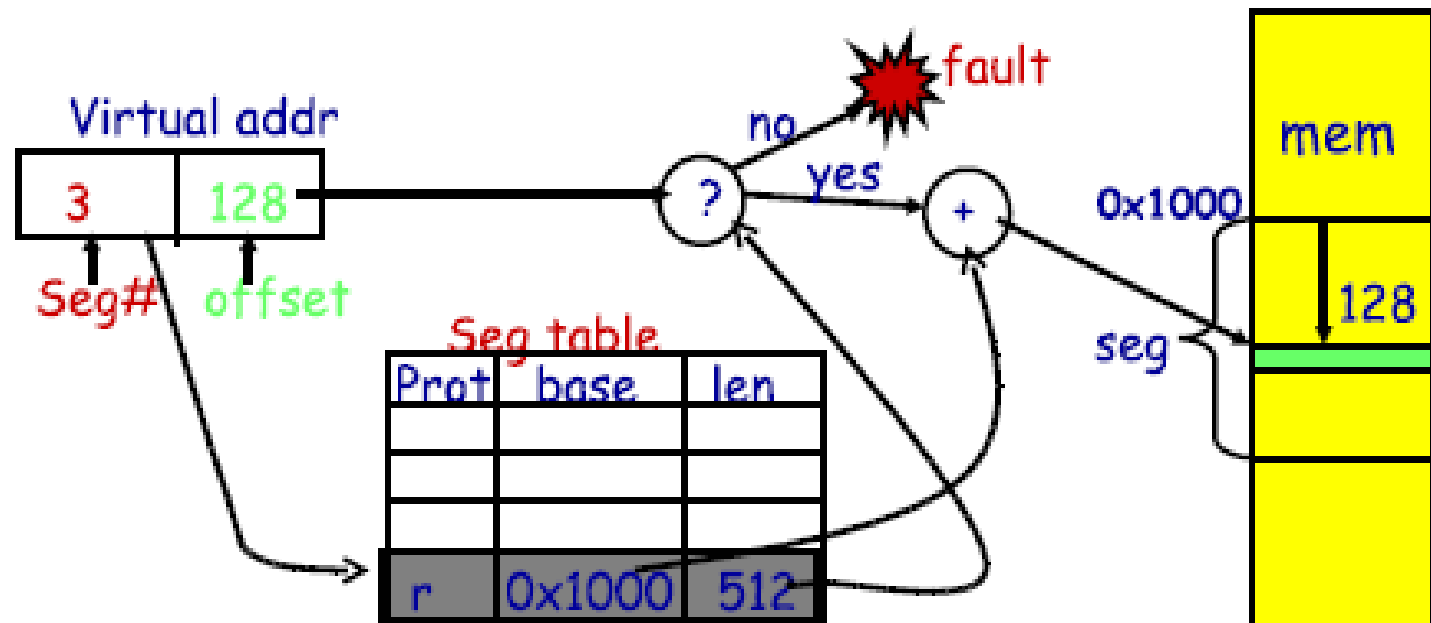
- Non contiguous allocation
 - Split a program in different non contiguous segments of variable size
- Let processes have many base/bound registers
 - Address space built from many segments
 - Can share/protect memory at segment granularity
- Must specify segment as part of virtual address

Segmentation



Segmentation mechanism


- Each process has a segment table
 - Each virtual address indicates a segment and offset:
 - Top bits of addr select segment, low bits select offset



Segmentation example

- 4-bit segment number (1st digit), 12 bit offset (last 3)
 - Where is 0x0240? 0x1108? 0x265c? 0x3002? 0x1600?

Seg	base	bounds	rw
0	0x4000	0x6fff	10
1	0x0000	0x4fff	11
2	0x3000	0xffff	11
3			00

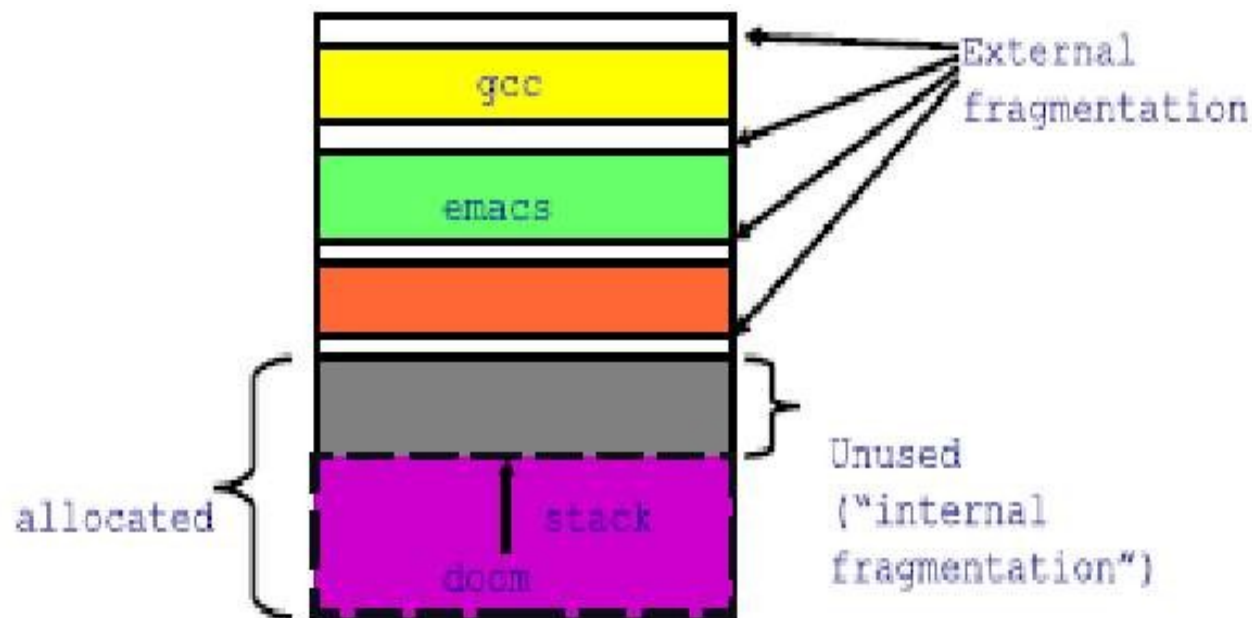


Segmentation tradeoffs

- Advantages
 - Multiple segments per process
 - Allows sharing
- Disadvantages
 - N byte segment needs N contiguous bytes of physical memory
 - Fragmentation (need moving segments)

Remember fragmentation problem

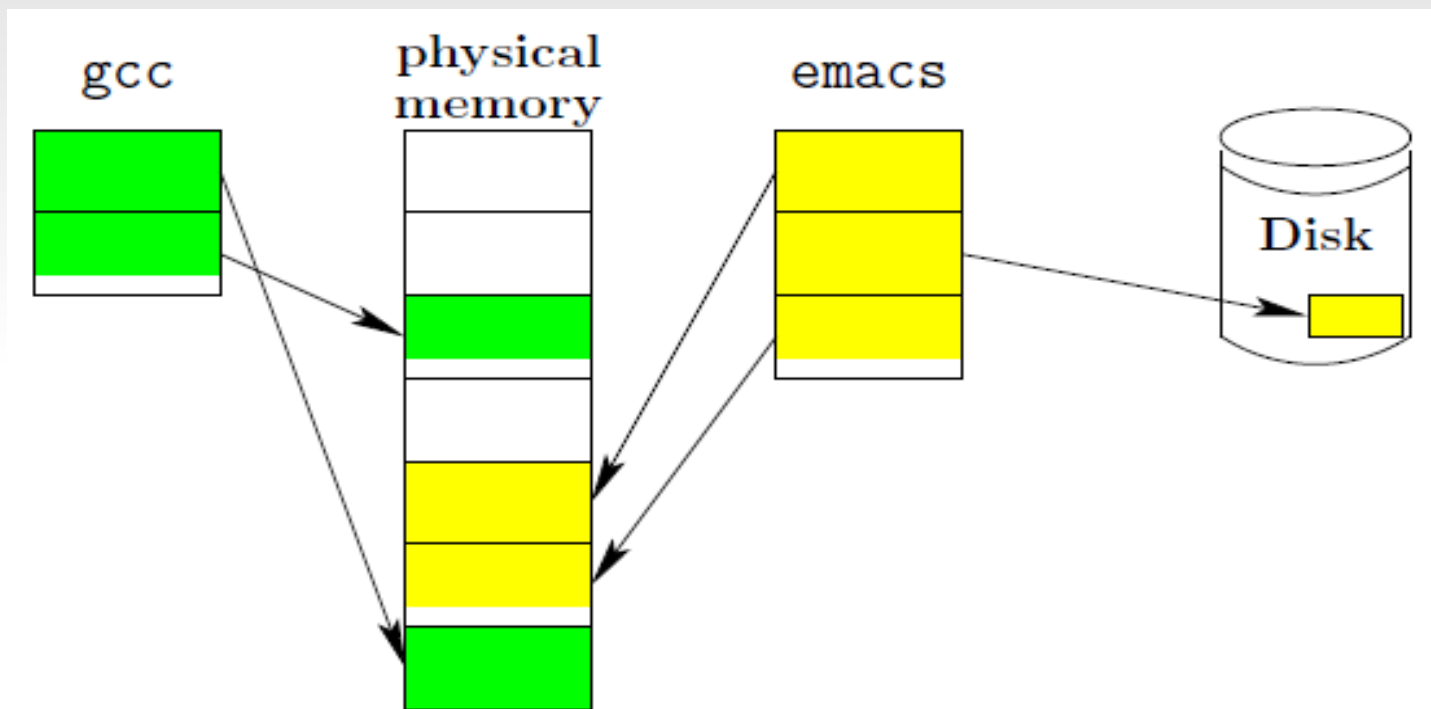
- Fragmentation => inability to use free memory
- Overtime:
 - Variable-size pieces = many small holes (external fragmentation)



Paging

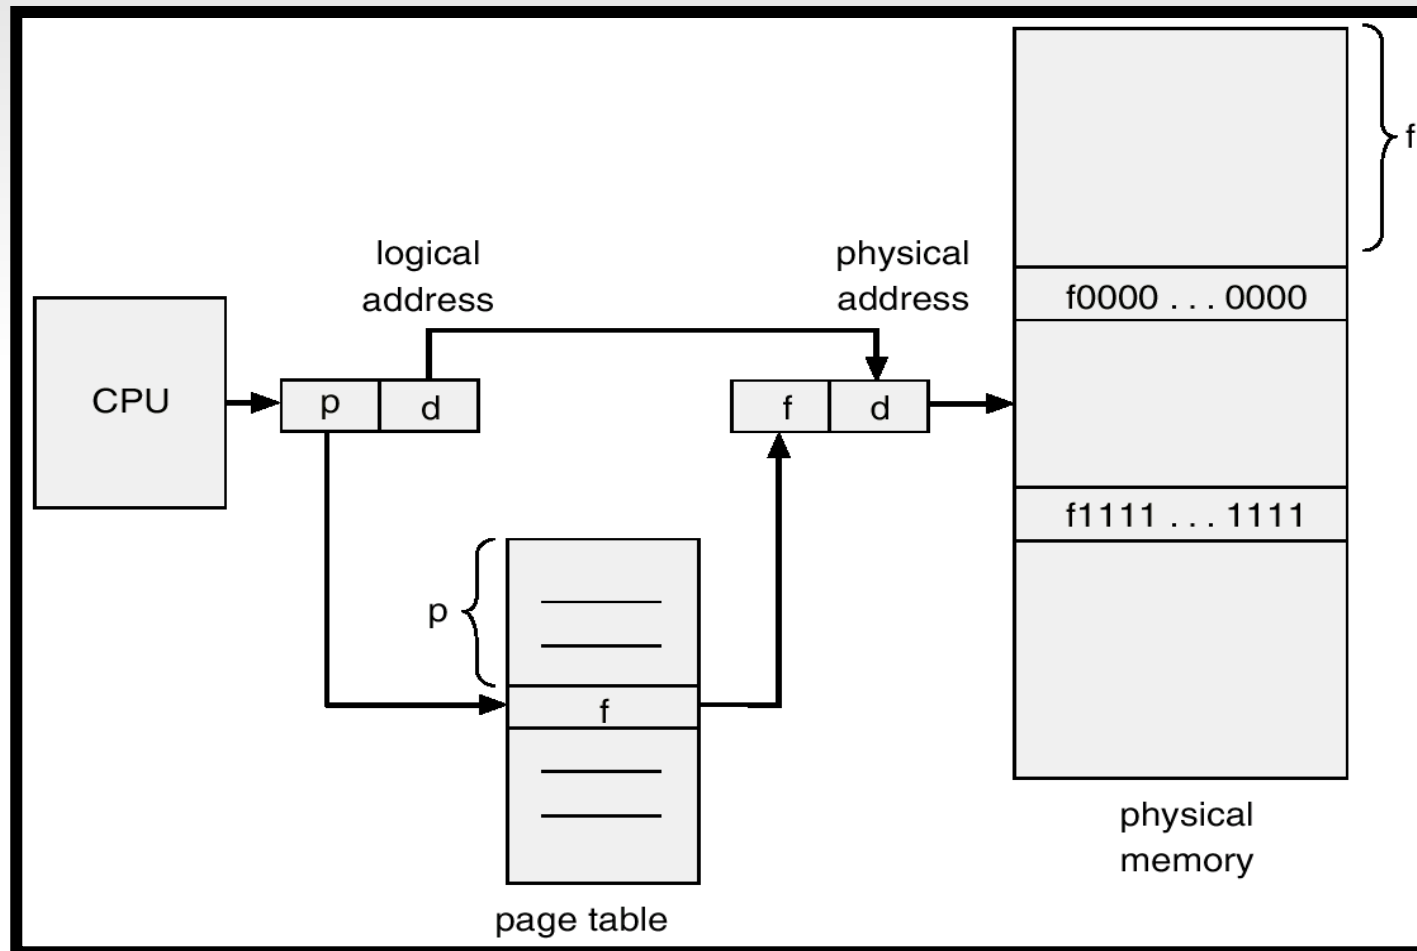
- Virtual memory is divided into small pages
 - Pages are fixed size
 - A page is contiguous
- Map virtual pages to physical block
 - Non contiguous allocation of blocks
 - Each process has a separate mapping but can share the same physical block
 - MMU
- OS gains control on certain operations
 - Non allocated pages trap to OS on access
 - Read only pages trap to OS on write
 - OS can change the mapping

Paging



- Page table
 - Global or per process

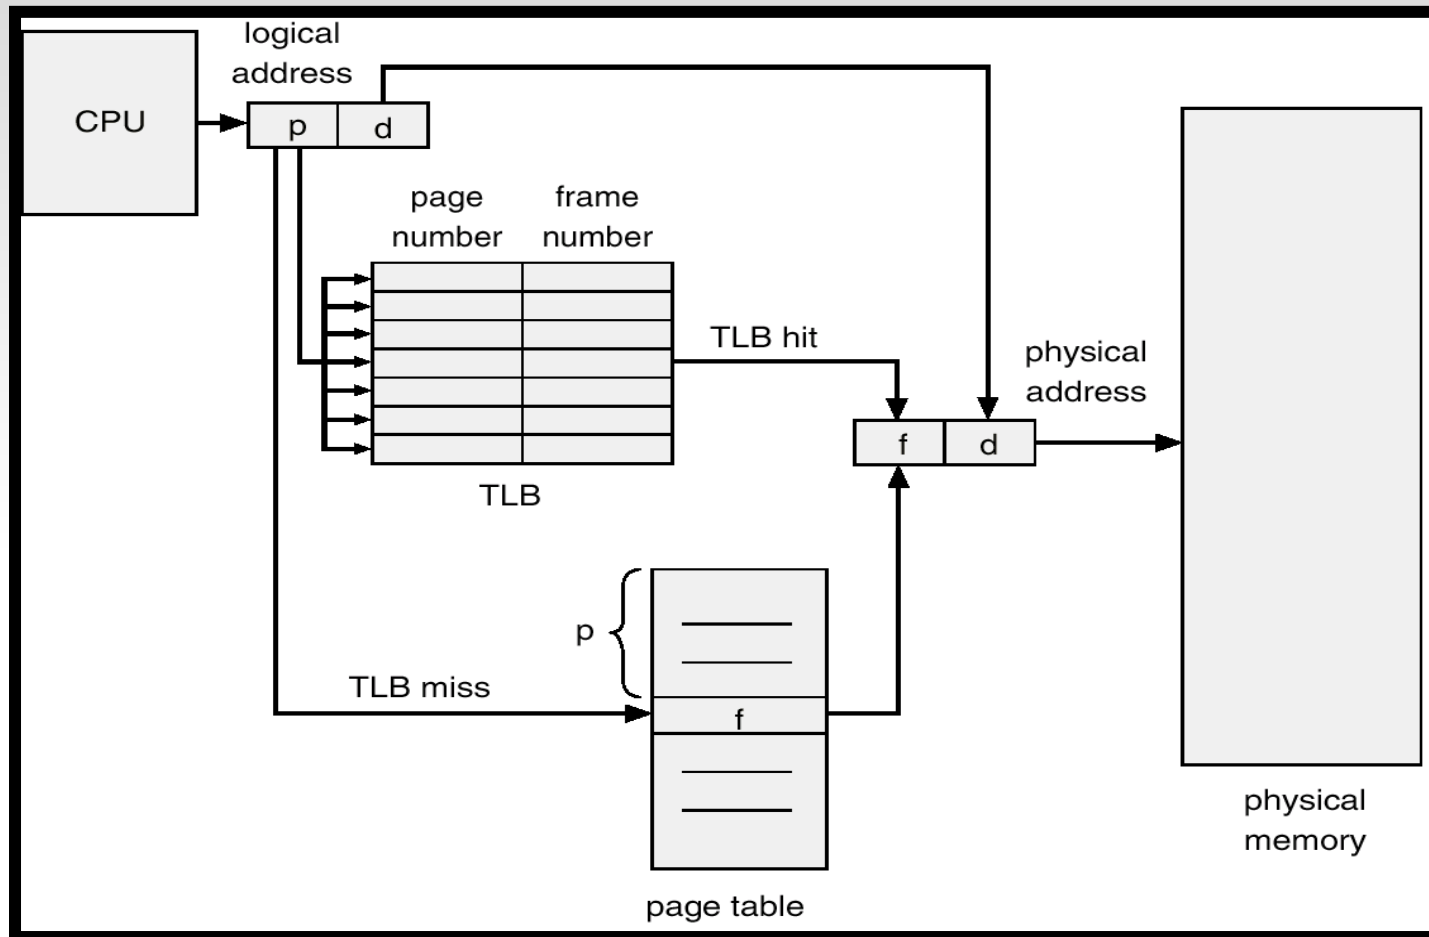
Virtual address translation



Problem : translation speed

- Require extra memory references on each load/store
 - Cache recently used translations
 - Locality principle
 - High probability that the next required address is close
- Translation Lookaside Buffer (TLB)
 - Fast (small) associative memory which can perform a parallel search
 - Typical TLB
 - Hit time : 1 clock cycle
 - Miss rate 1%
 - TLB management : hardware or software

TLB



- What to do when switch address space ?
 - Flush the TLB
 - Tag each entry with the process's id
- Update TLB on page fault (add/remove TLB entries)

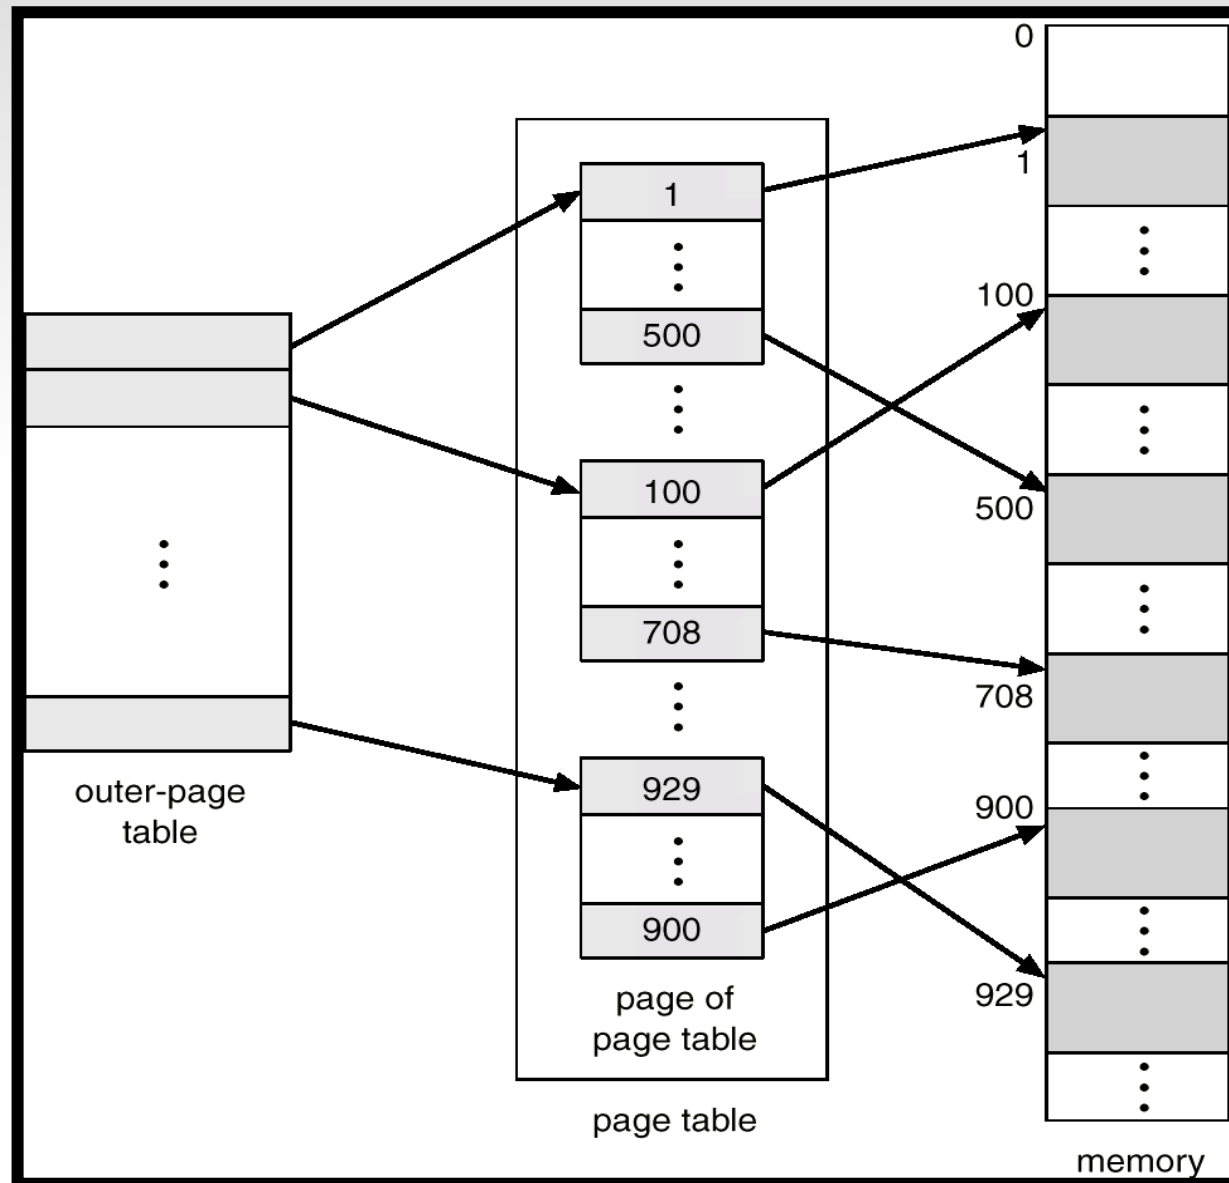
Problem : page table size

- Flat page tables are huge
- Example
 - 4GB of virtual memory (32 bits address)
 - 4KB pages
 - 20bits page number, 12 bits offset
 - 1MB page table size :<

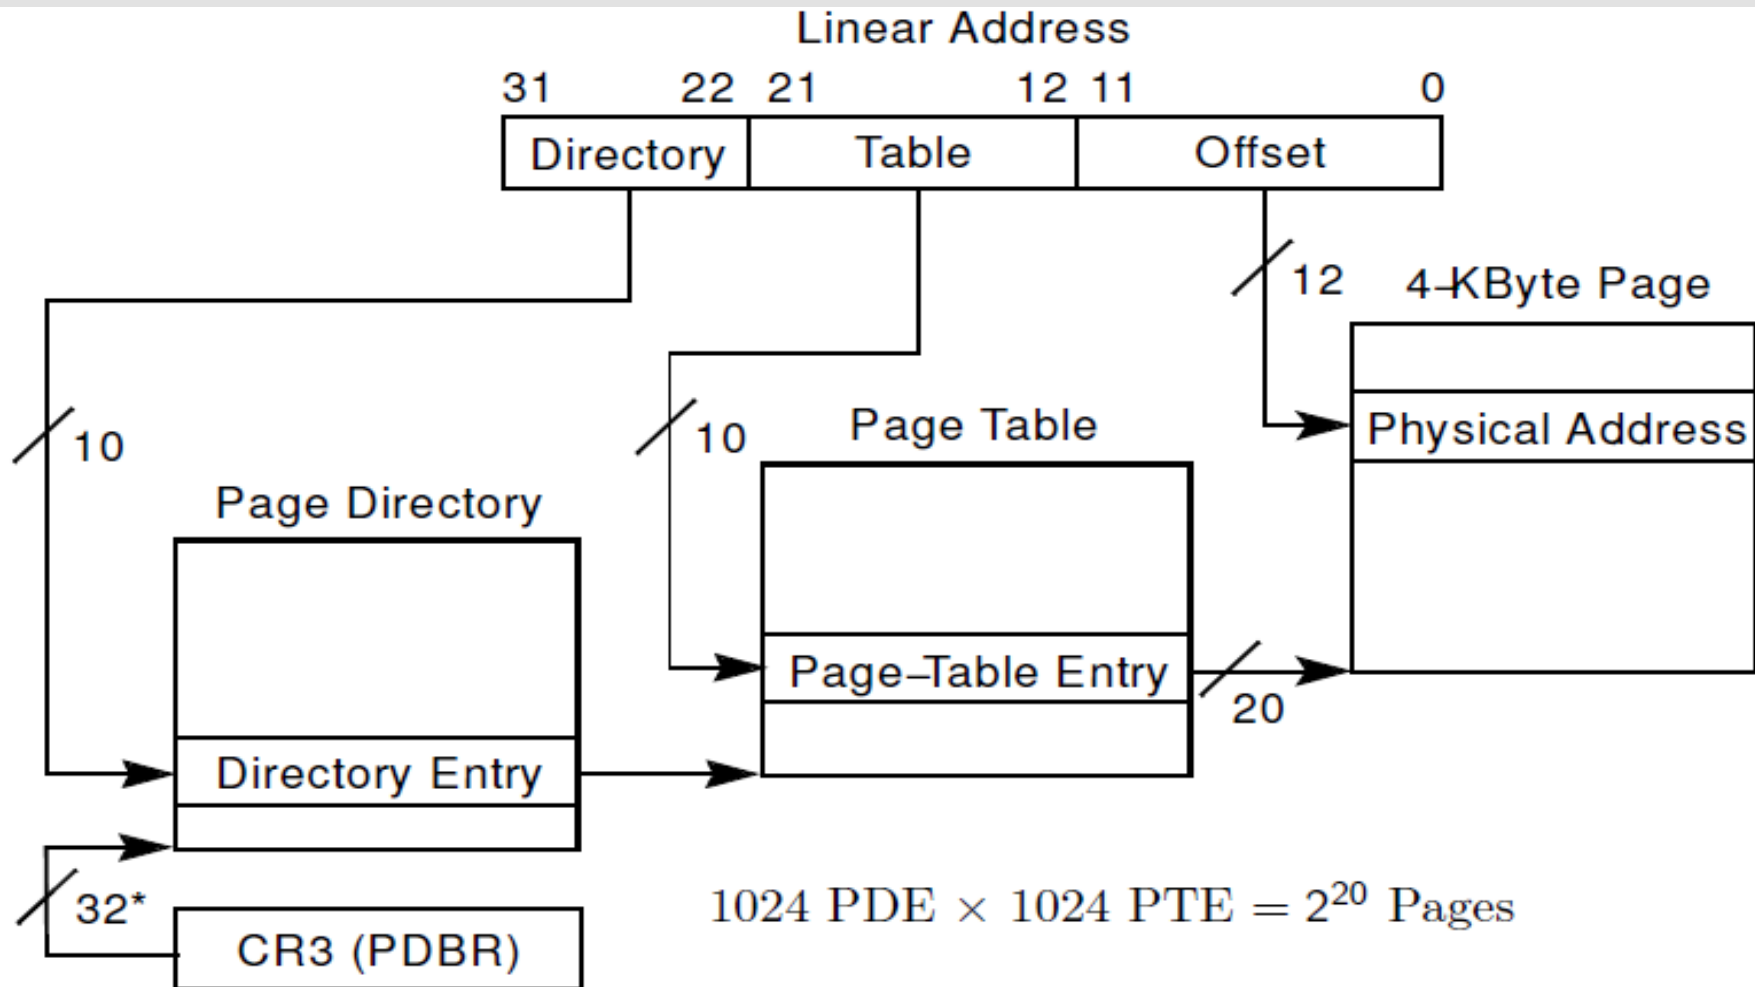
Multi-level page tables

- Reduce the size of page tables in memory
- Structured page tables in 2 or more levels
 - All the page tables are not present in memory all the time
 - Some page tables are stored on disk and fetched if necessary
- Based on a on-demand paging mechanism

Example: two level pages



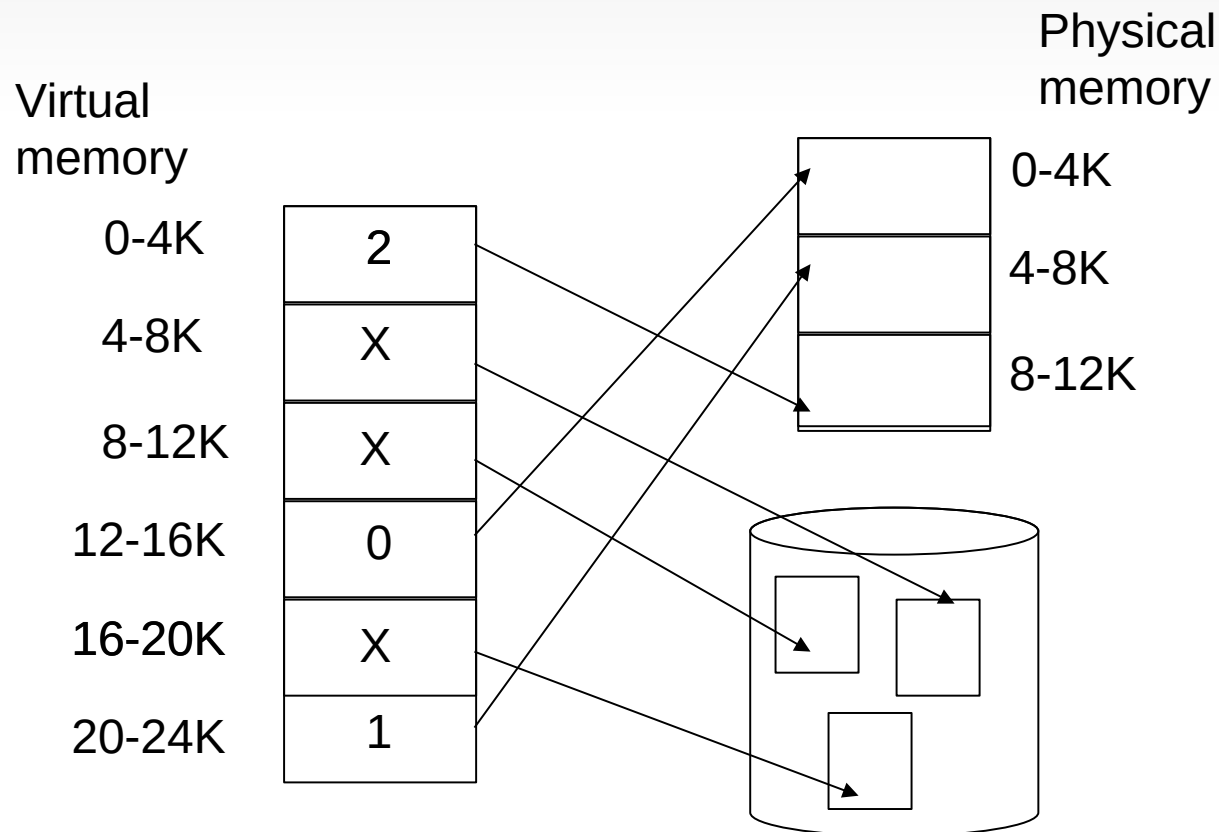
Example: two level pages



*32 bits aligned onto a 4-KByte boundary

On demand paging

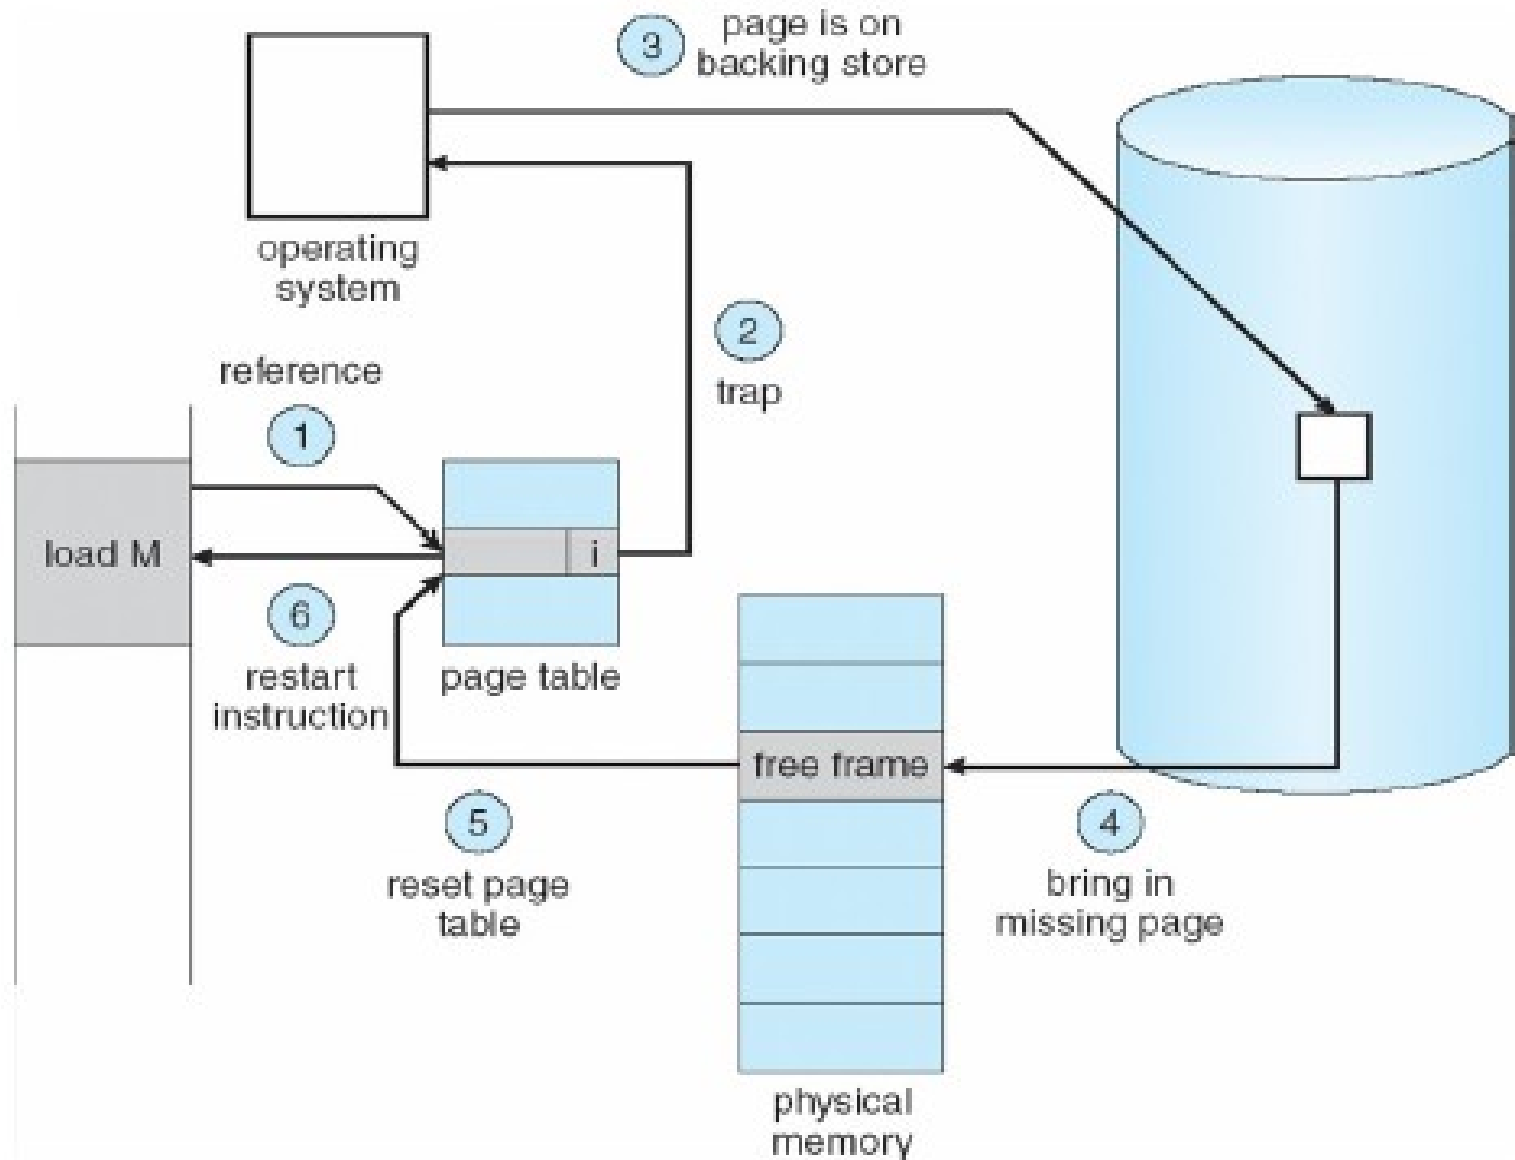
- Virtual memory > physical memory
 - Some pages are not present in memory (X)
 - Stored on disk



Page fault

- Access to an absent page
 - Presence bit
 - Page fault (Trap to OS)
- Page fault management
 - Find a free physical frame
 - If there is a free frame; use it
 - Else, select a page to replace (to free a frame)
 - Save the replaced page on disk if necessary (dirty page)
 - Load the page from disk in the physical frame
 - Update page table
 - Restart instruction
- Require a presence bit, a dirty bit, a disk @ in the page table
- Different page replacement algorithms

On demand paging

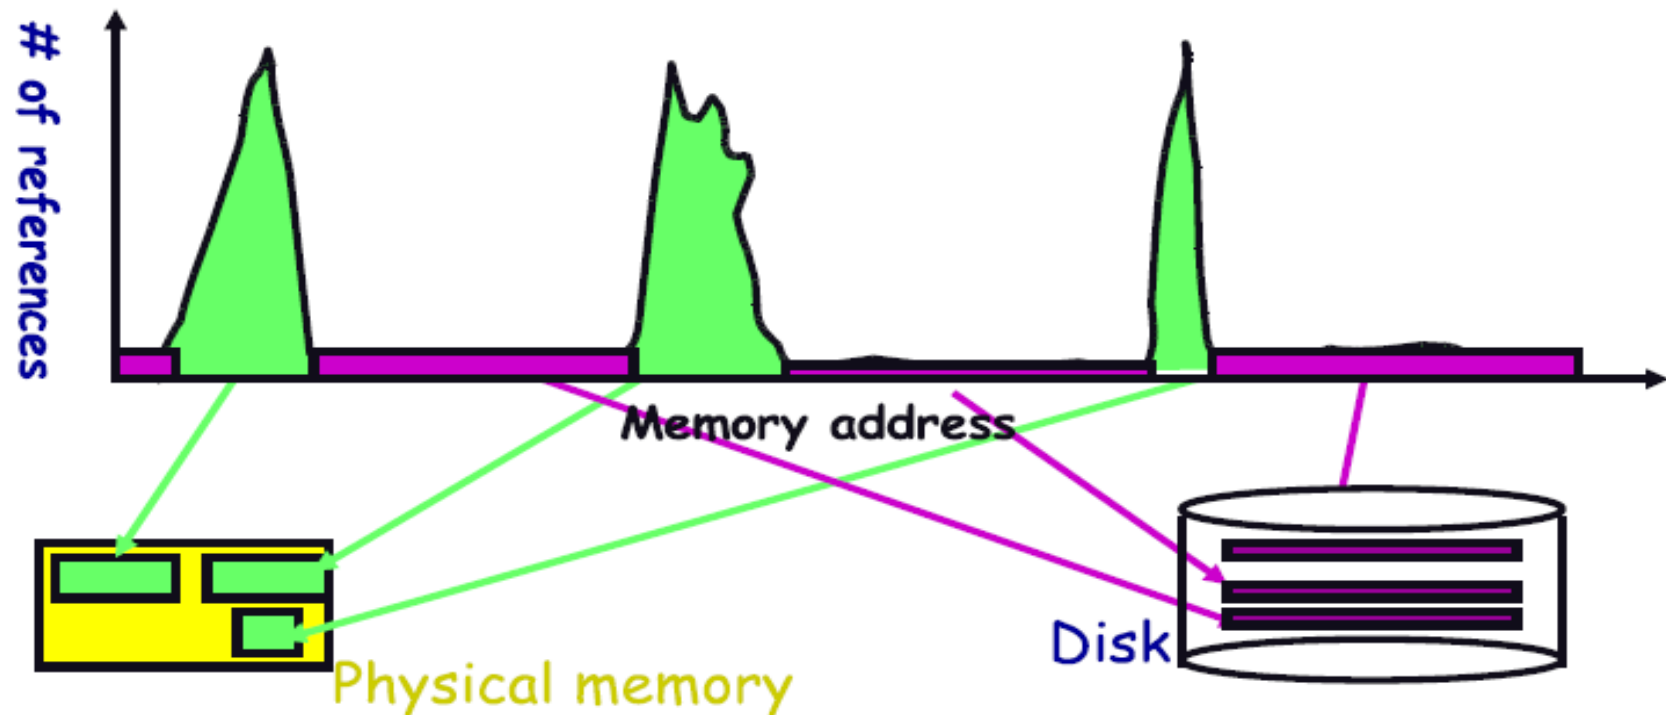


Page replacement algorithms

- Working set model
- Algorithms
 - Optimal
 - FIFO
 - Second chance
 - LRU

Working set model

- Disk much much slower than memory (RAM)
 - Goal: run at memory (not disk) speed
- 90/10 rule: 10% of memory gets 90% of memory refs
 - So, keep that 10% in real memory, the other 90% on disk



Optimal page replacement

- What is optimal (if you knew the future)?
 - Replace pages that will not be used for longest period of time
- Example
 - Reference string : 0,1,2,3,0,1,4,0,1,2,3,4,1,2
 - 4 physical frames:

0	0	3
1	1	1
2	2	2
3	4	4

6 pages faults

FiFo

- Evict oldest page in system
- Example
 - Reference string : 0,1,2,3,0,1,4,0,1,2,3,4,1,2
 - 4 physical frames:

10 page faults

0	4	4	4	4	3	3
1	1	0	0	0	0	4
2	2	2	1	1	1	1
3	3	3	3	2	2	2

- Implementation: just a list (updated on page fault)

LRU page replacement

- Approximate optimal with least recently used
 - Because past often predicts the future
- Example
 - Reference string : 0,1,2,3,0,1,4,0,1,2,3,4,1,2
 - 4 physical frames:

0	0	0	0	4
1	1	1	1	1
2	4	4	3	3
3	3	2	2	2

8 page faults

LRU implementation

- Expensive
 - Need specific hardware
 - Track access without page fault
- Approximate LRU
 - The aging algorithm
 - Add a counter for each page (the date)
 - On a page access, all page counters are shifted right, inject 1 for the accessed page, else 0
 - On a page replacement, remove the page with the lowest counter

Aging : example

Accessed page	Date Page0	Date Page1	Date Page2	Order pages /date
	000	000	000	
Page 0	1 00	000	000	P0,P1=P2
Page 1	0 1 0	1 00	000	P1,P0,P2
Page 2	00 1	0 1 0	1 00	P2,P1,P0
Page 1	000	1 0 1	0 1 0	P1,P2,P0

P0 is the oldest

Second chance

- Simple FIFO modification
 - Use an access bit R for each page
 - Set to 1 when page is referenced
 - Periodically reset by hardware
 - Inspect the R bit of the oldest page (of the FIFO list)
 - If 0 : replace the page
 - If 1 : clear the bit, put the page at the end of the list, and repeat
- Approximation of LRU
 - don't have to parse all pages

Page buffering

- Naïve paging
 - Page replacement : 2 disk IO per page fault
- Reduce the IO on the critical path
 - Keep a pool of free frames
 - Fetch the page in an already free page
 - Swap out a page in background

Paging

- Separate linking from memory concern
- Simplifies allocation, free and swap
- Eliminate external fragmentation
- May leverage internal fragmentation

Resources you can read

- http://en.wikipedia.org/wiki/Page_table
 - Wikipedia can always be useful
- Operating System Concepts, 10th Edition, Abraham Silberschatz, Peter B. Galvin, Greg Gagne
 - <http://os-book.com/>
 - Chapters 9 & 10
- Modern Operating Systems, Andrew Tanenbaum
 - <http://www.cs.vu.nl/~ast/books/mos2/>
 - Chapter 4